



# Kubernetes Foundations Training

Tomasz Cholewa <[tomasz@cloudowski.com](mailto:tomasz@cloudowski.com)>

# Agenda

## Day 1

Container basics

Operating docker containers

Running containers from available images

Providing configuration and storage for containers

Multi-container configurations using docker-compose

Building and publishing container images

# Agenda

## Day 2

Kubernetes architecture and basic components

Managing objects in k8s

Pod and its main features

Scaling of applications runnins as pods

Storing data on persistent storage

# Agenda

## Day 3

Exposing application using Service object

Exposing http applications using Ingress

Basic logs maintenance and monitoring

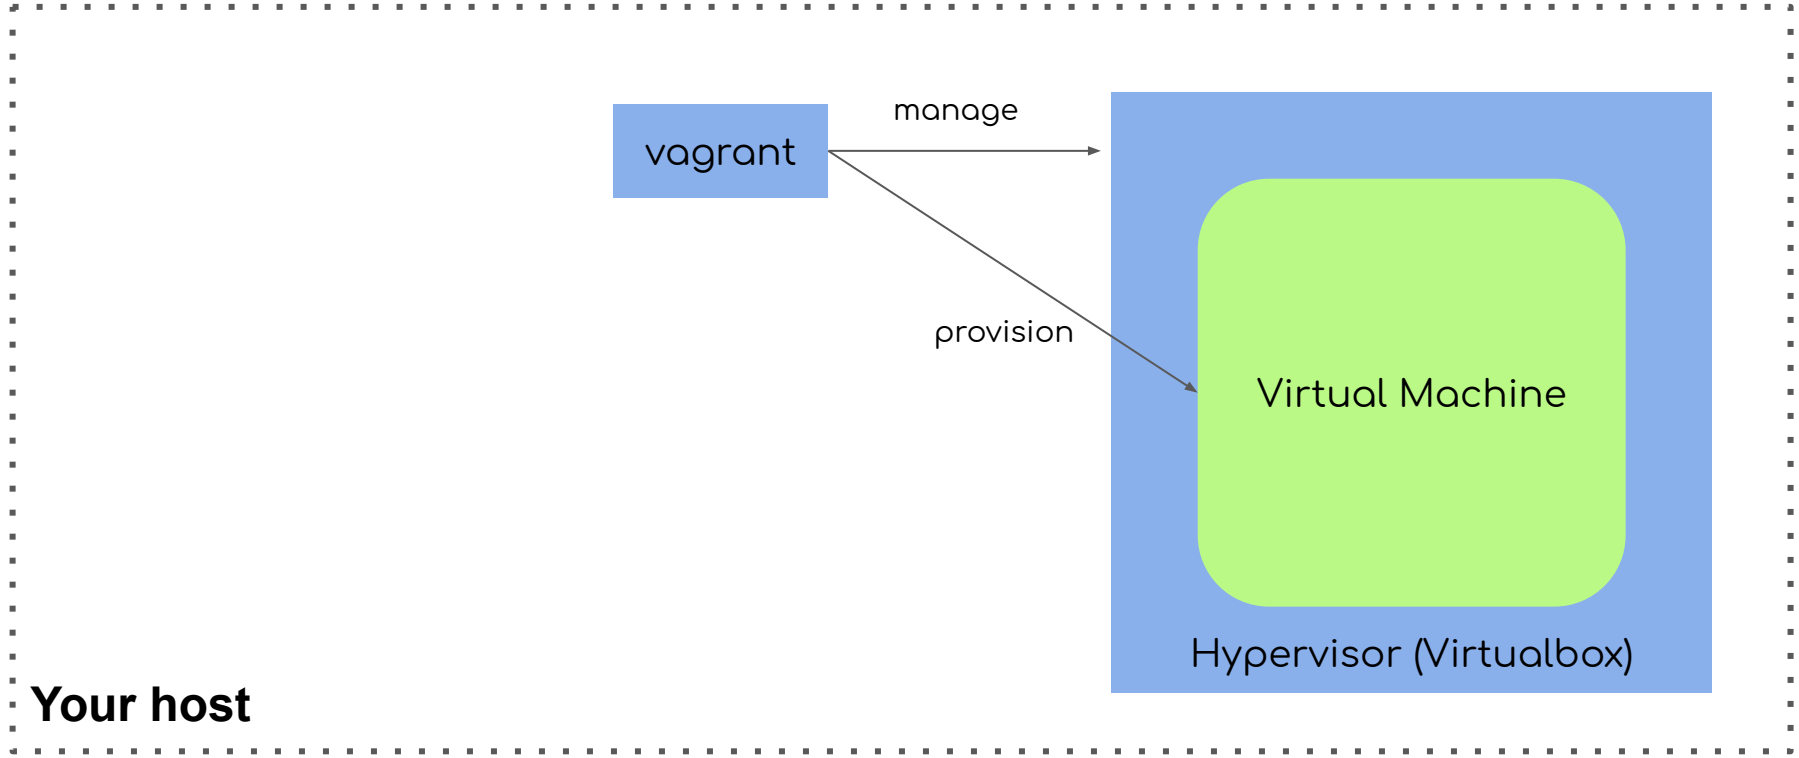
Managing applications using Deployment

Installing apps using Helm and Charts

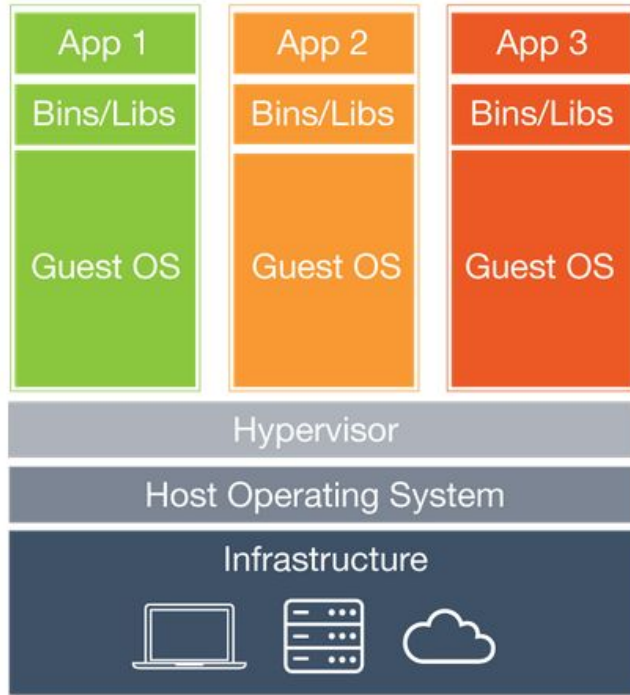
# Quest



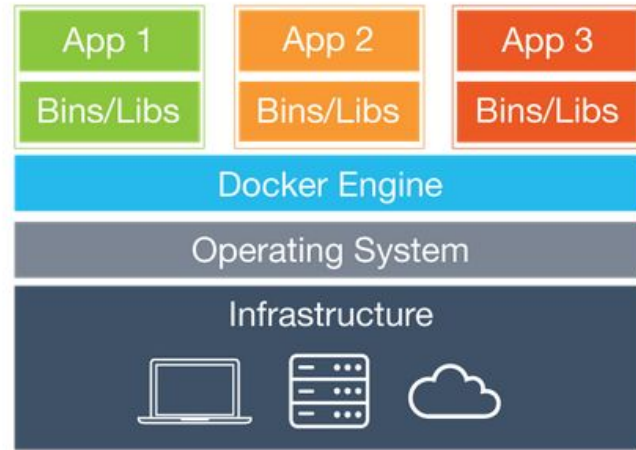
# Vagrant



# Containers vs. Virtual Machines




Virtual Machines



Containers

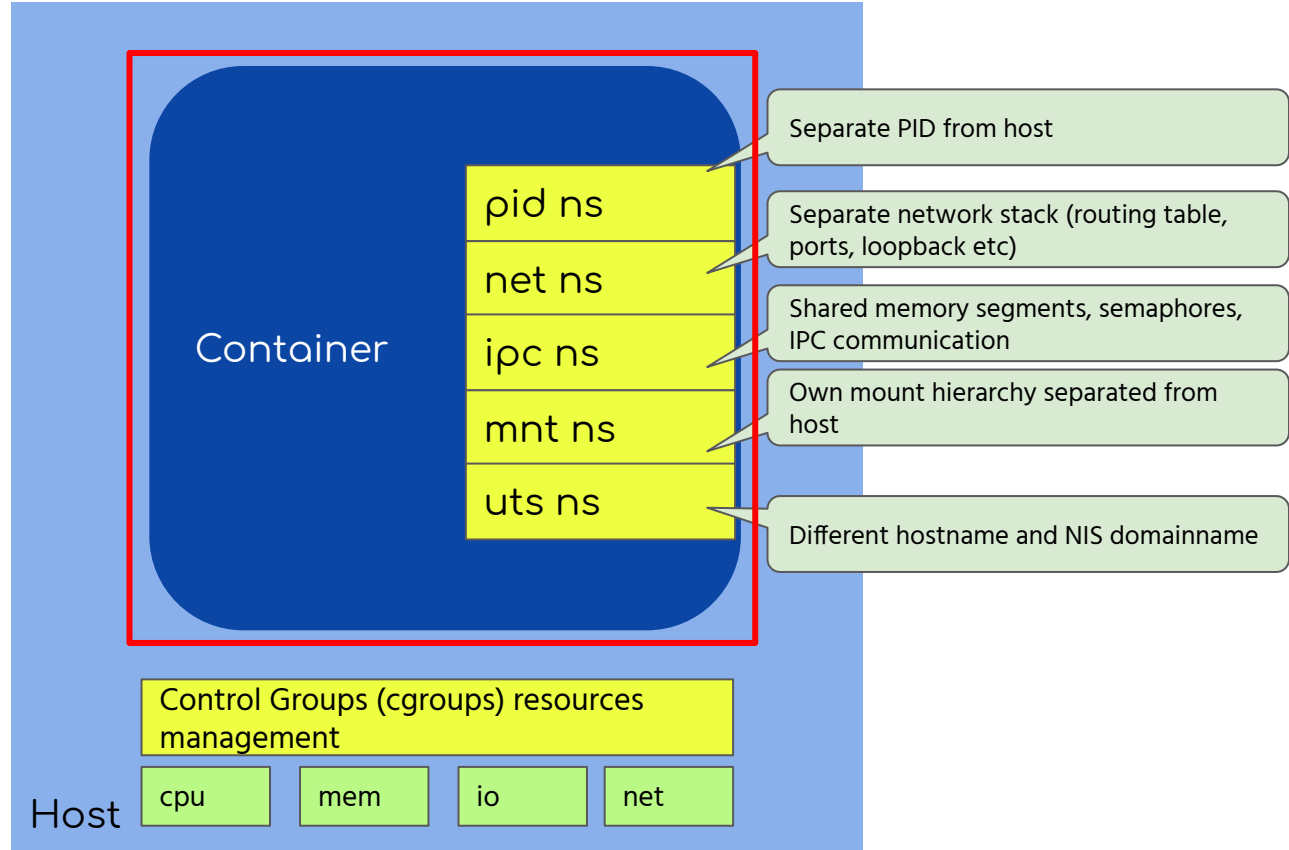
# Why containers?

1. Easy to use
2. Portability
3. Security
4. Speed






# Container isolation

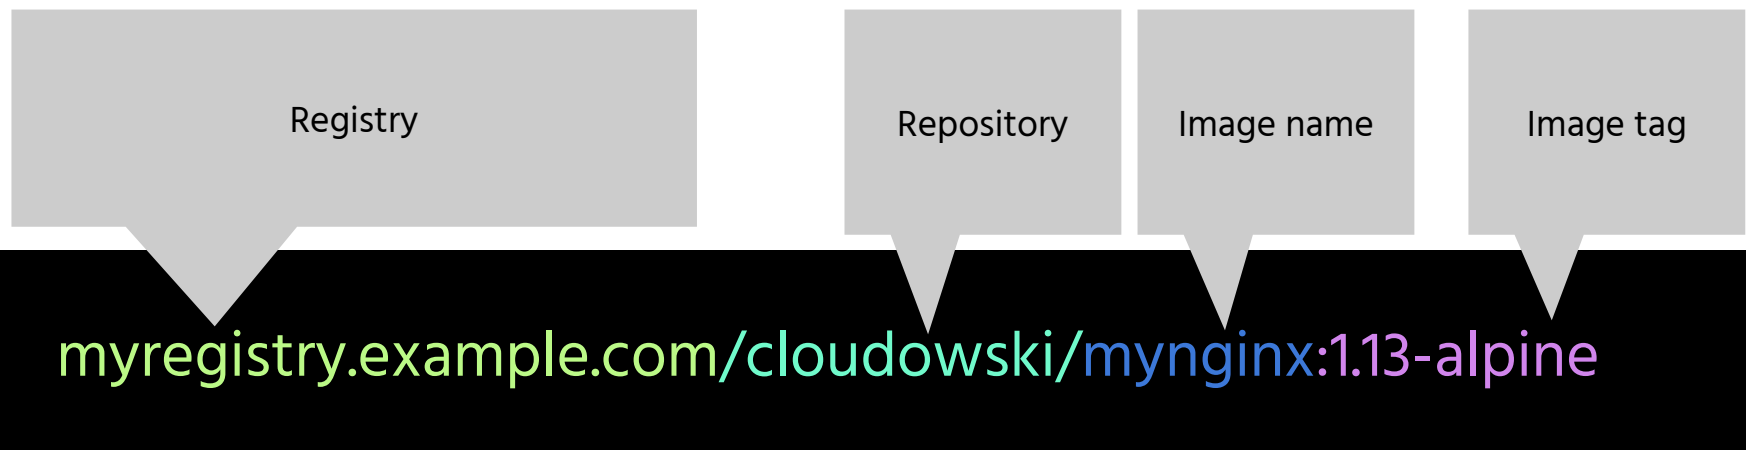


# Containers, peristent data and configuration



time


# Container image naming




# Image layers



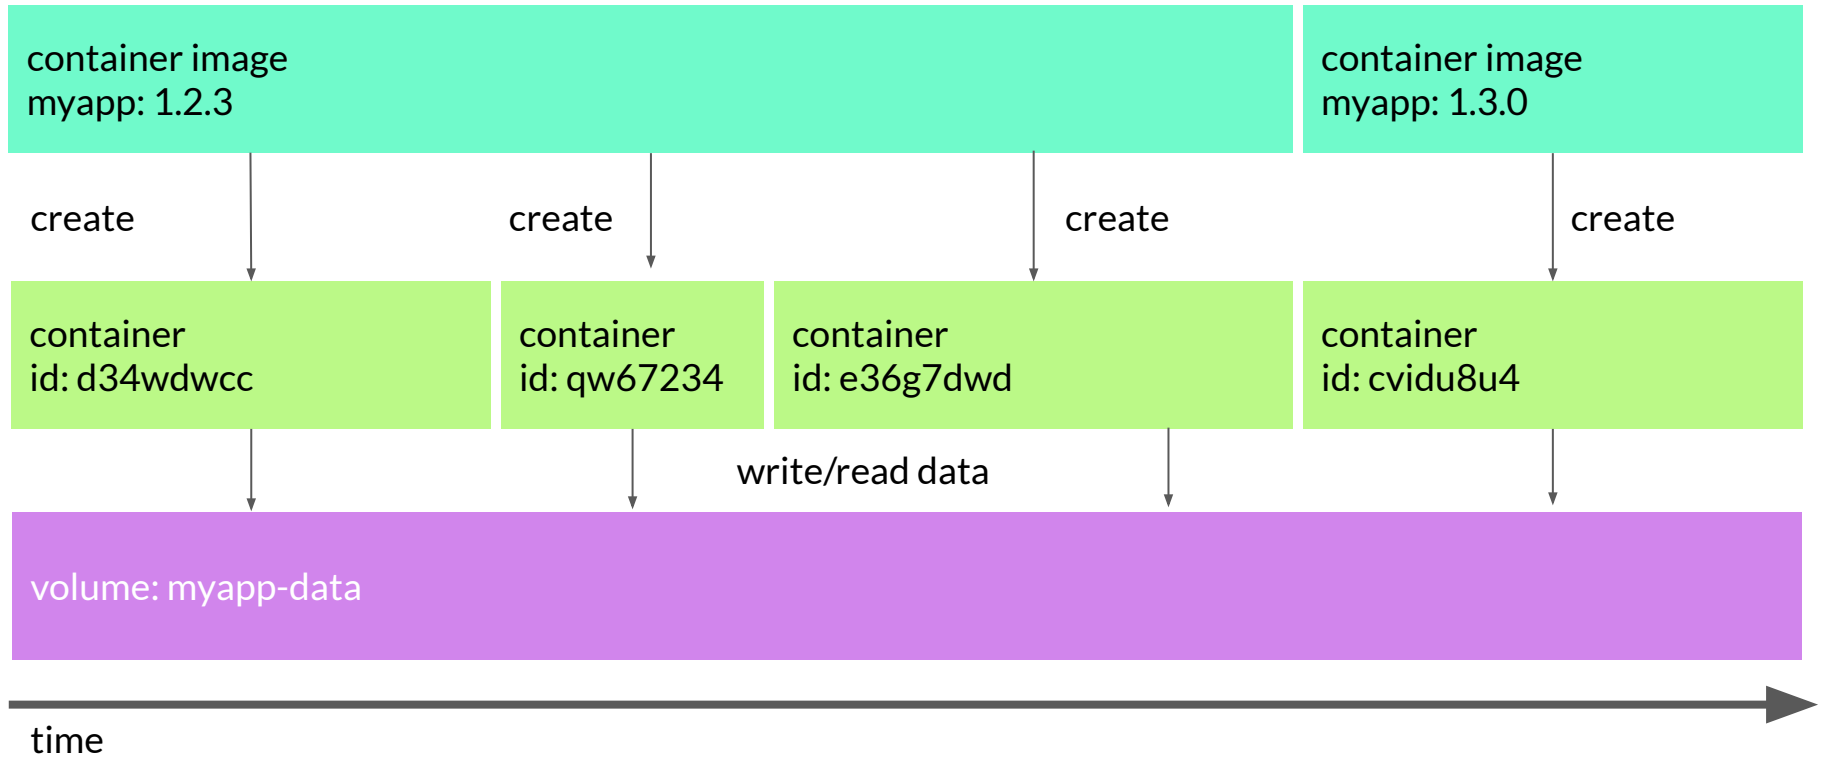
# Creating new image from container



# Traditional upgrade and configuration



# Ephemeral containers with persistent data



# 10 tips for building good container images

1. Keep them small



# 10 tips for building good container images

2. Single app/single purpose for single image

# 10 tips for building good container images

## 3. Leverage docker caching

# 10 tips for building good container images

4. Create few layers with only necessary tools

# 10 tips for building good container images

5. Always put EXPOSE for all exposed ports

# 10 tips for building good container images

## 6. Assign proper tags

# 10 tips for building good container images

7. Create universal Dockerfiles  
configurable for future releases

# 10 tips for building good container images

## 8. Handle signals properly

# 10 tips for building good container images

9. Always prefer official images instead of custom or built in-house



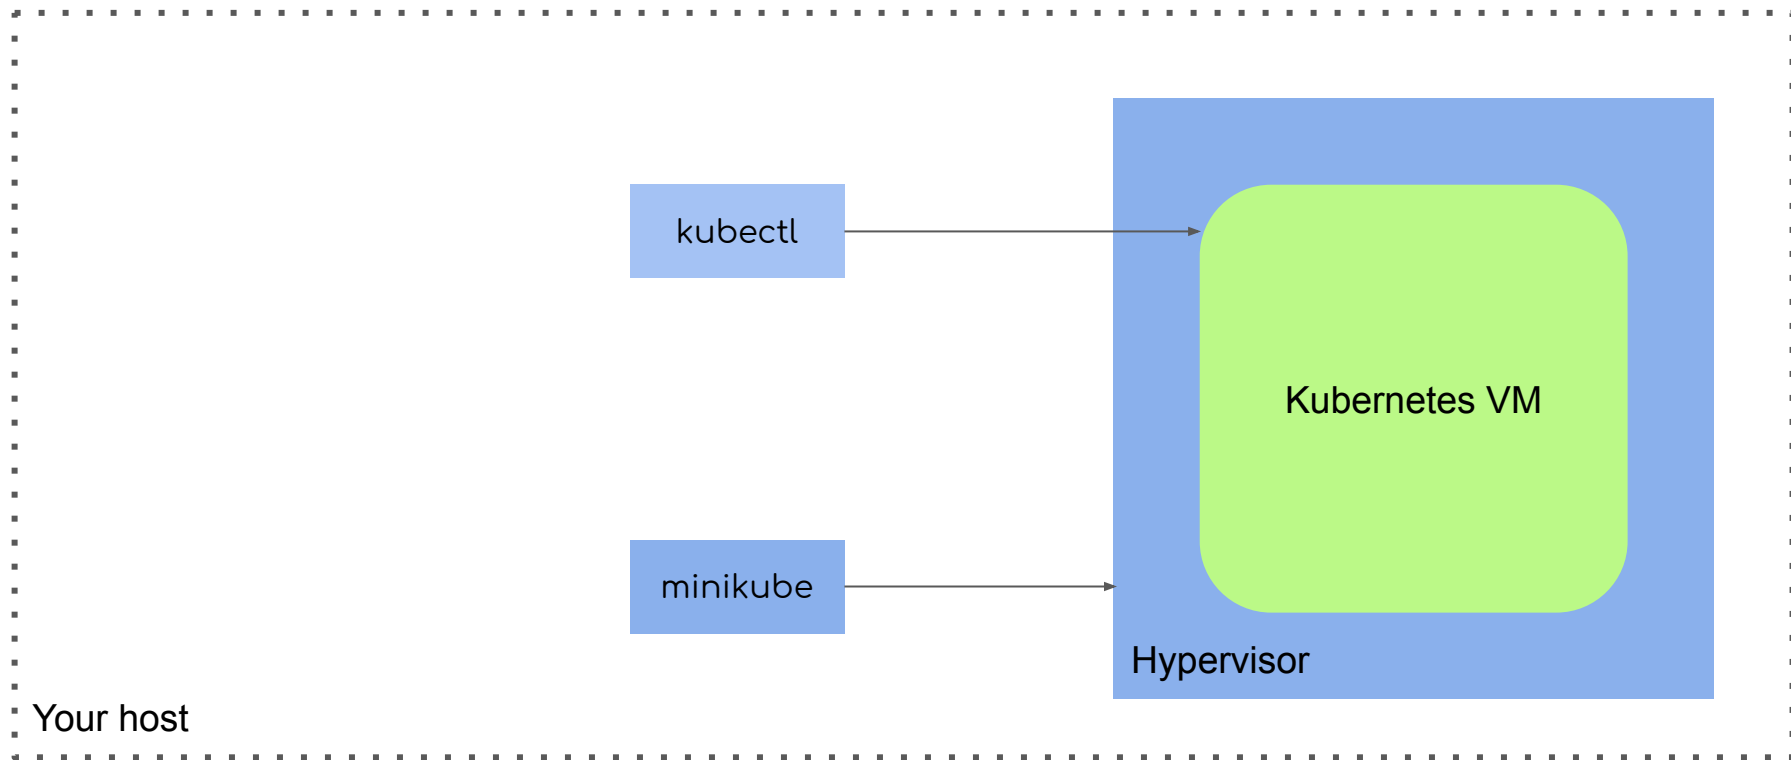
# 10 tips for building good container images

## 10. Automate build process using CI/CD




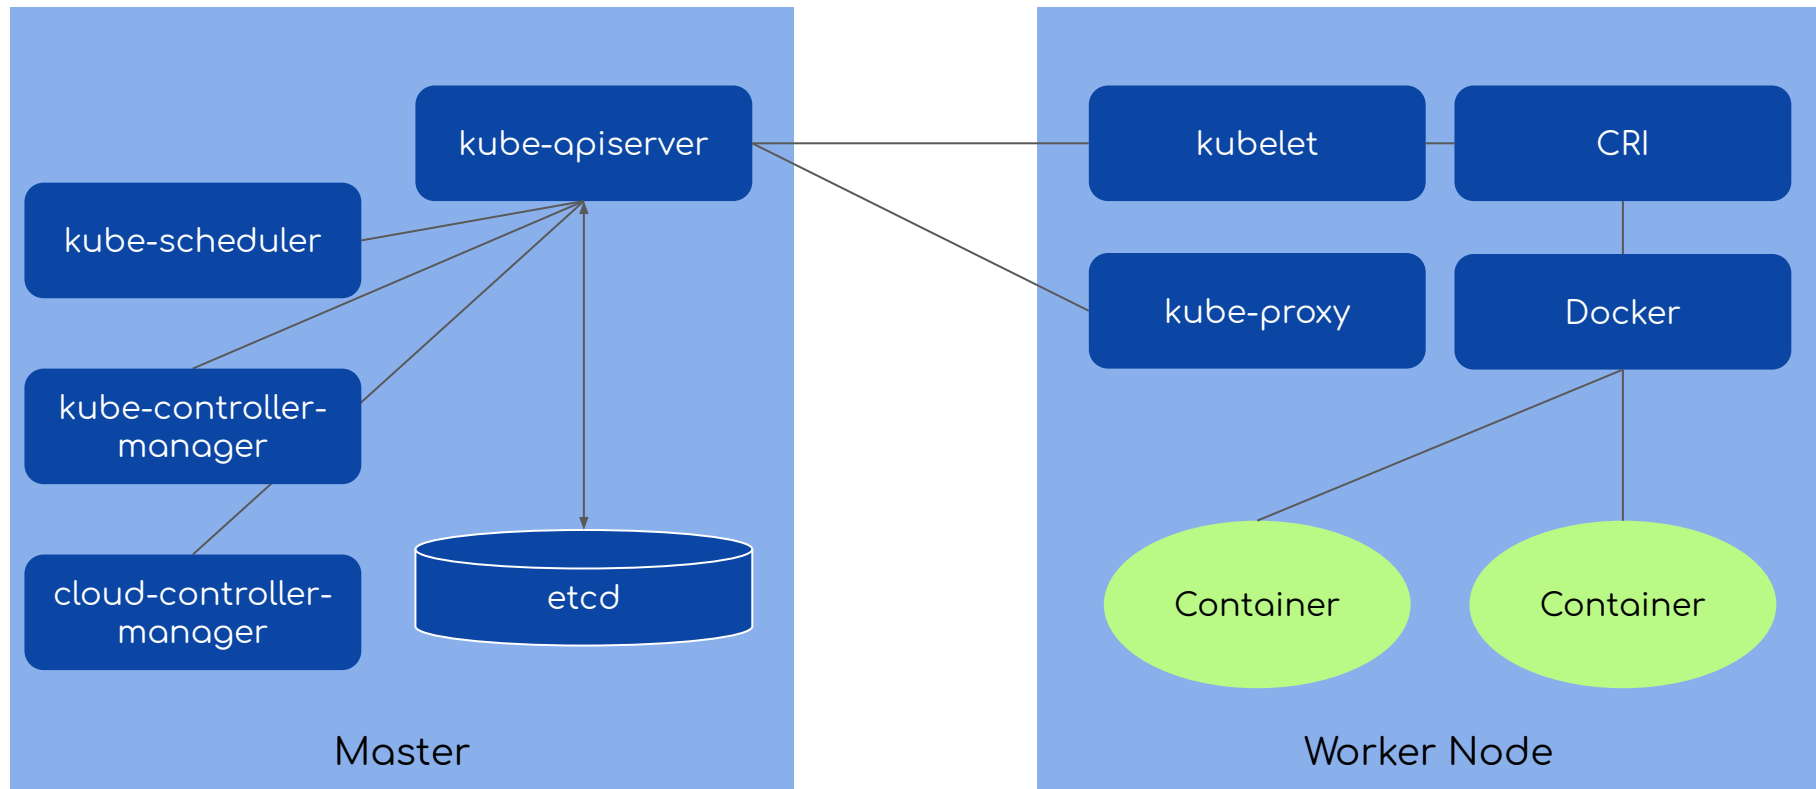
**KUBERNETES**

# Minikube overview



# What is Kubernetes?





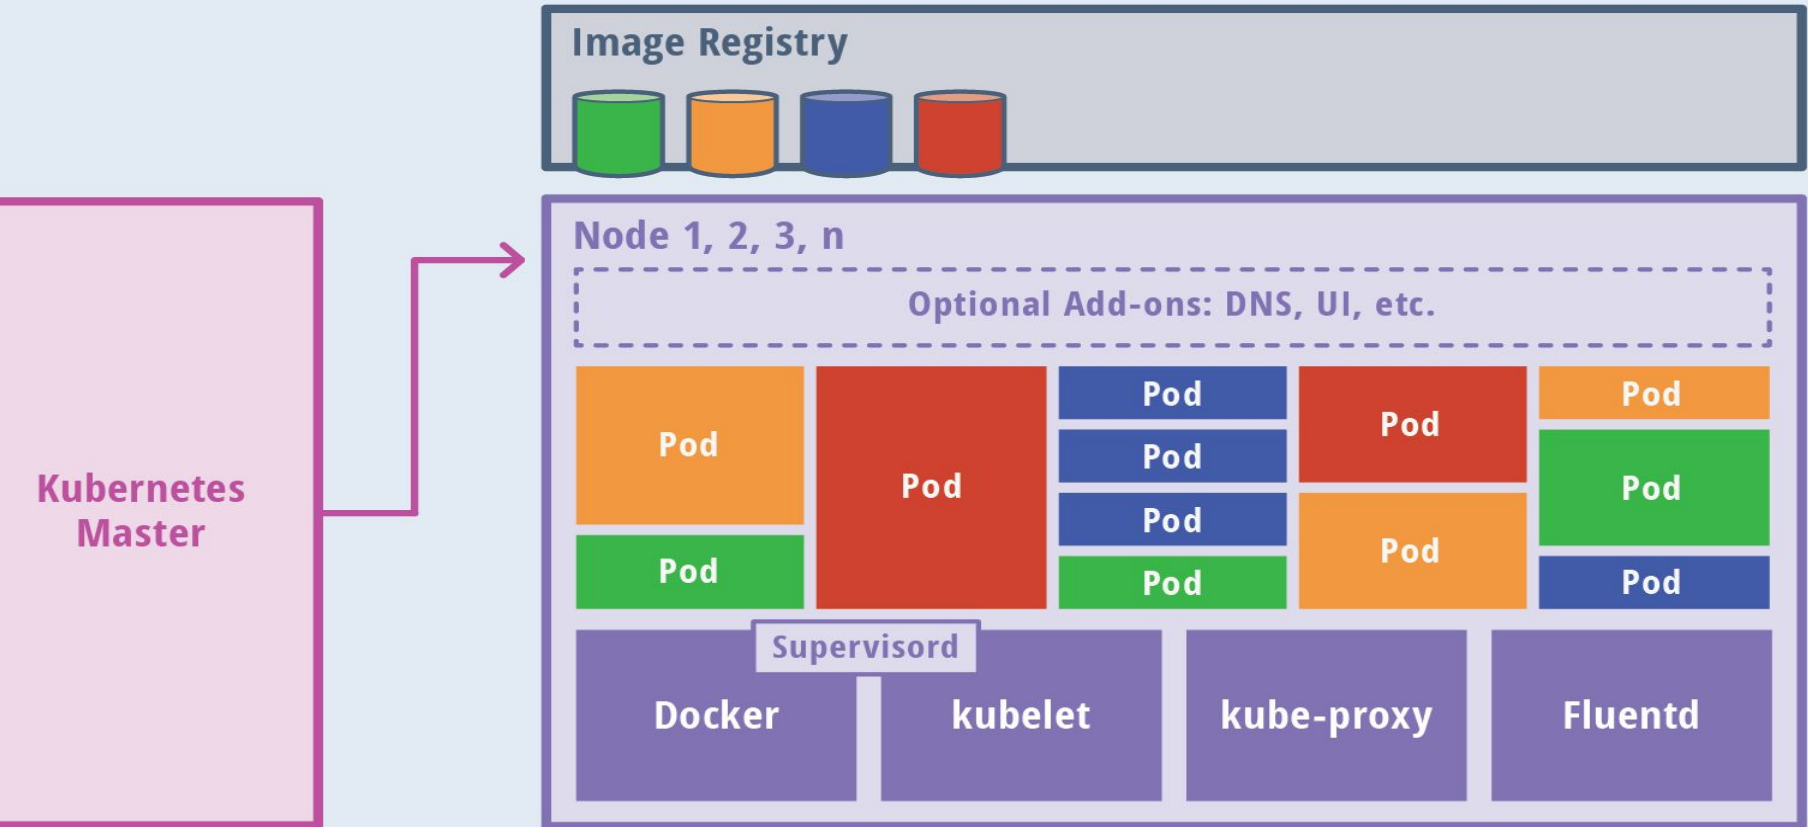
Kubernetes architecture details

# Why Kubernetes?

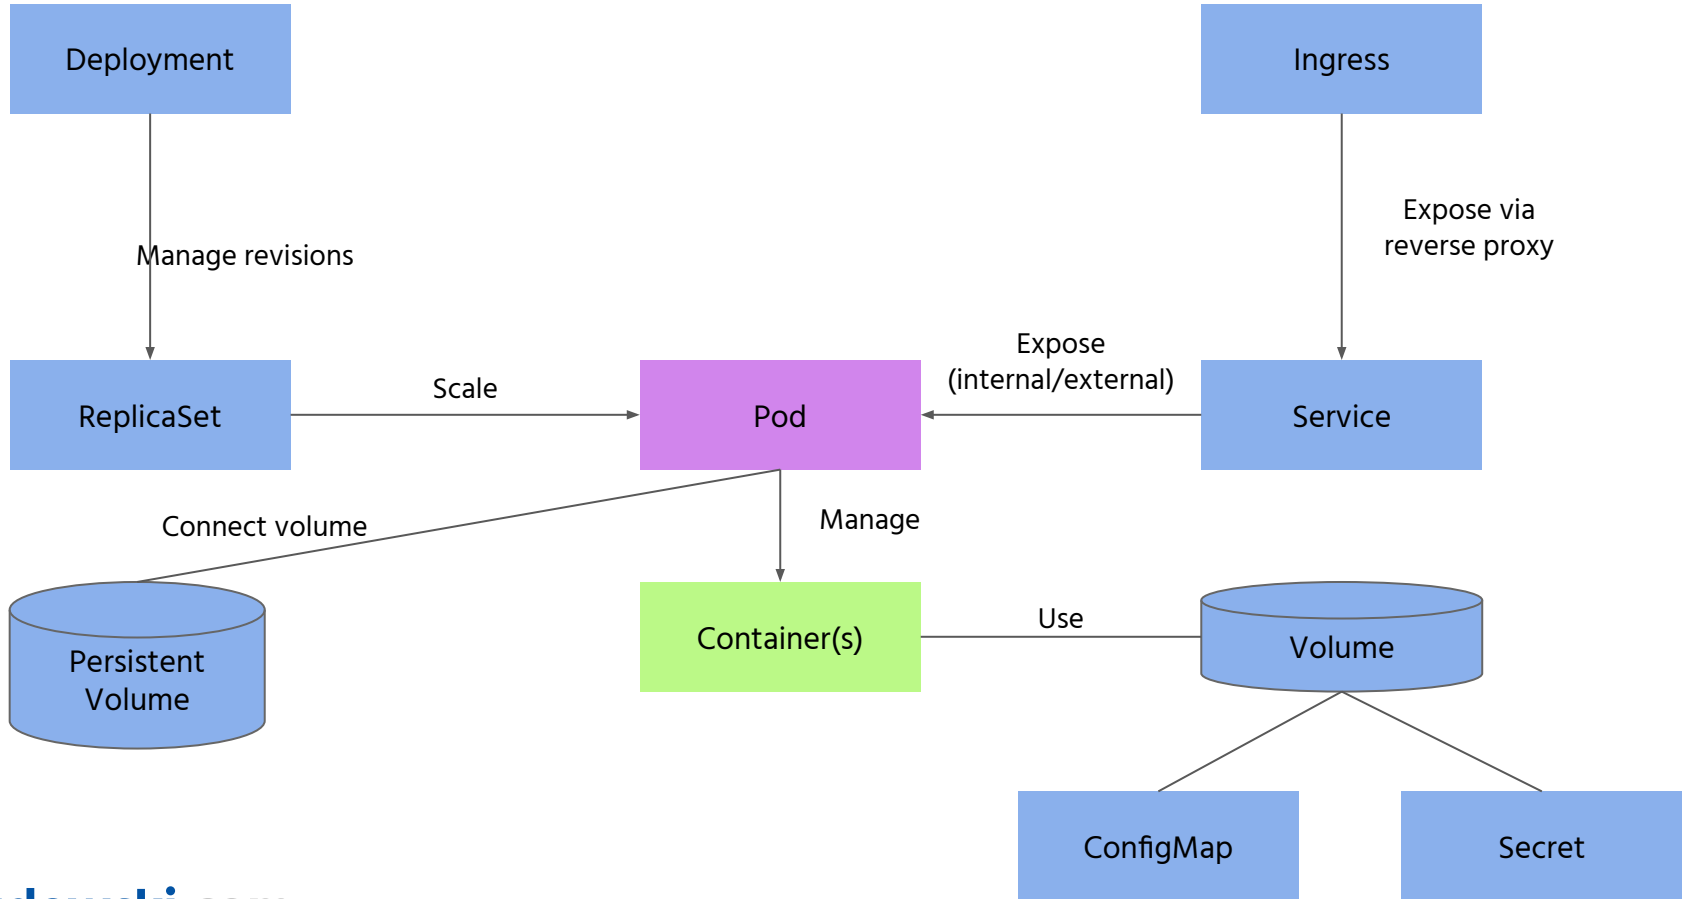
1. Portability
2. Security
- 3. Scalability**
4. Speed
- 5. Automation**
- 6. High Availability**
- 7. Extensibility**



# Kubernetes Node

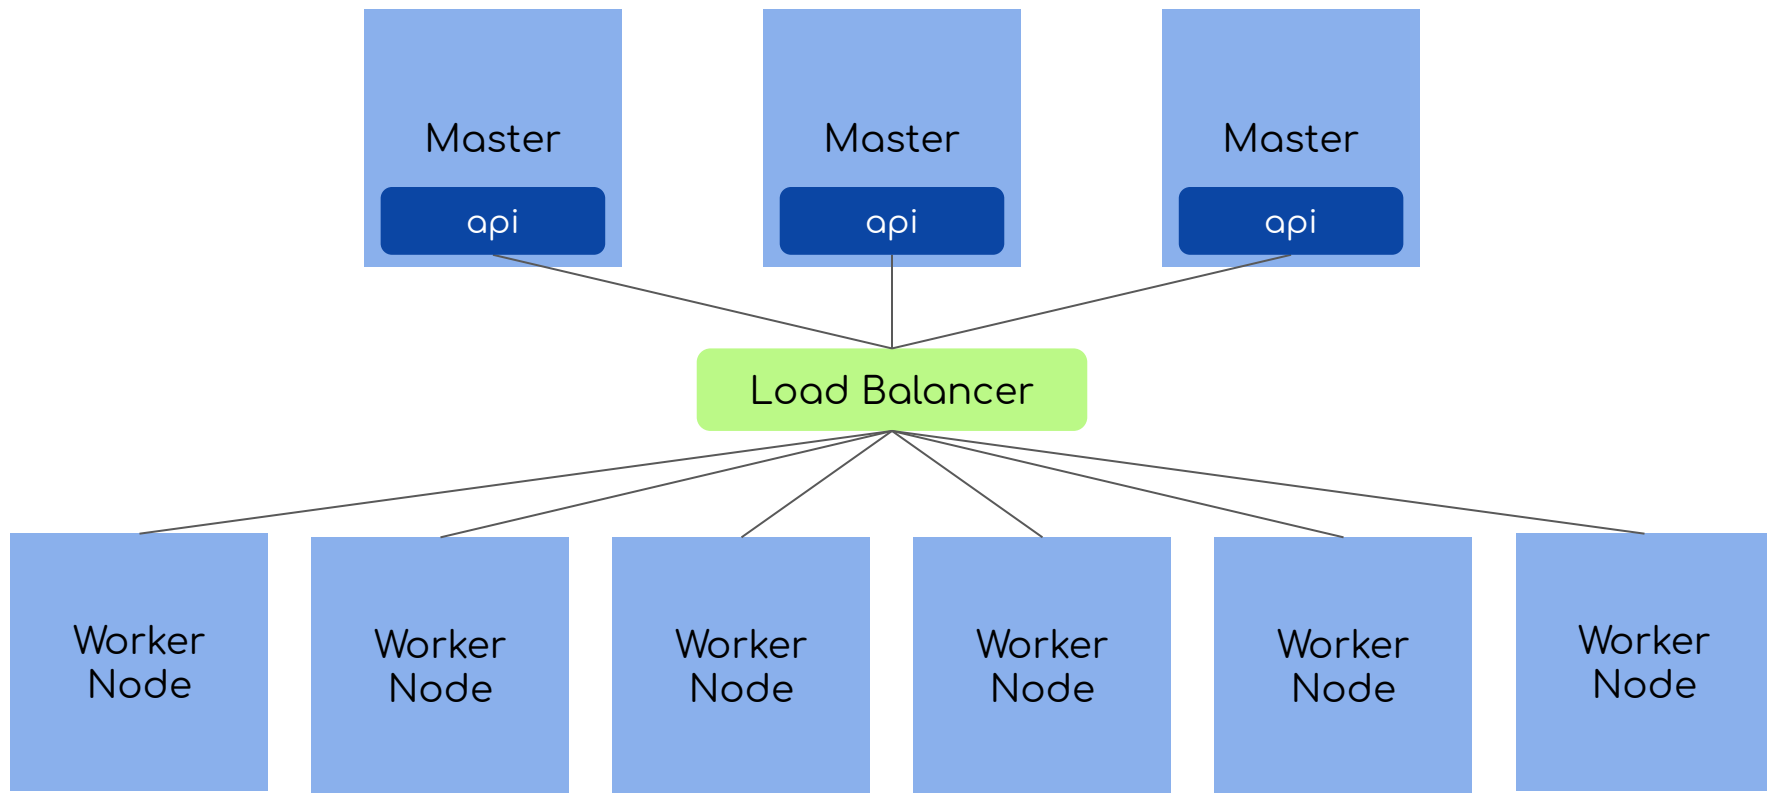


# Kubernetes map





# Kubernetes networking




Kubernetes architecture overview

# Rules of Kubernetes networking model

1. All containers can communicate with all other containers without NAT
2. All nodes can communicate with all containers (and vice-versa) without NAT
3. The IP that a container sees itself as is the same IP that others see it as

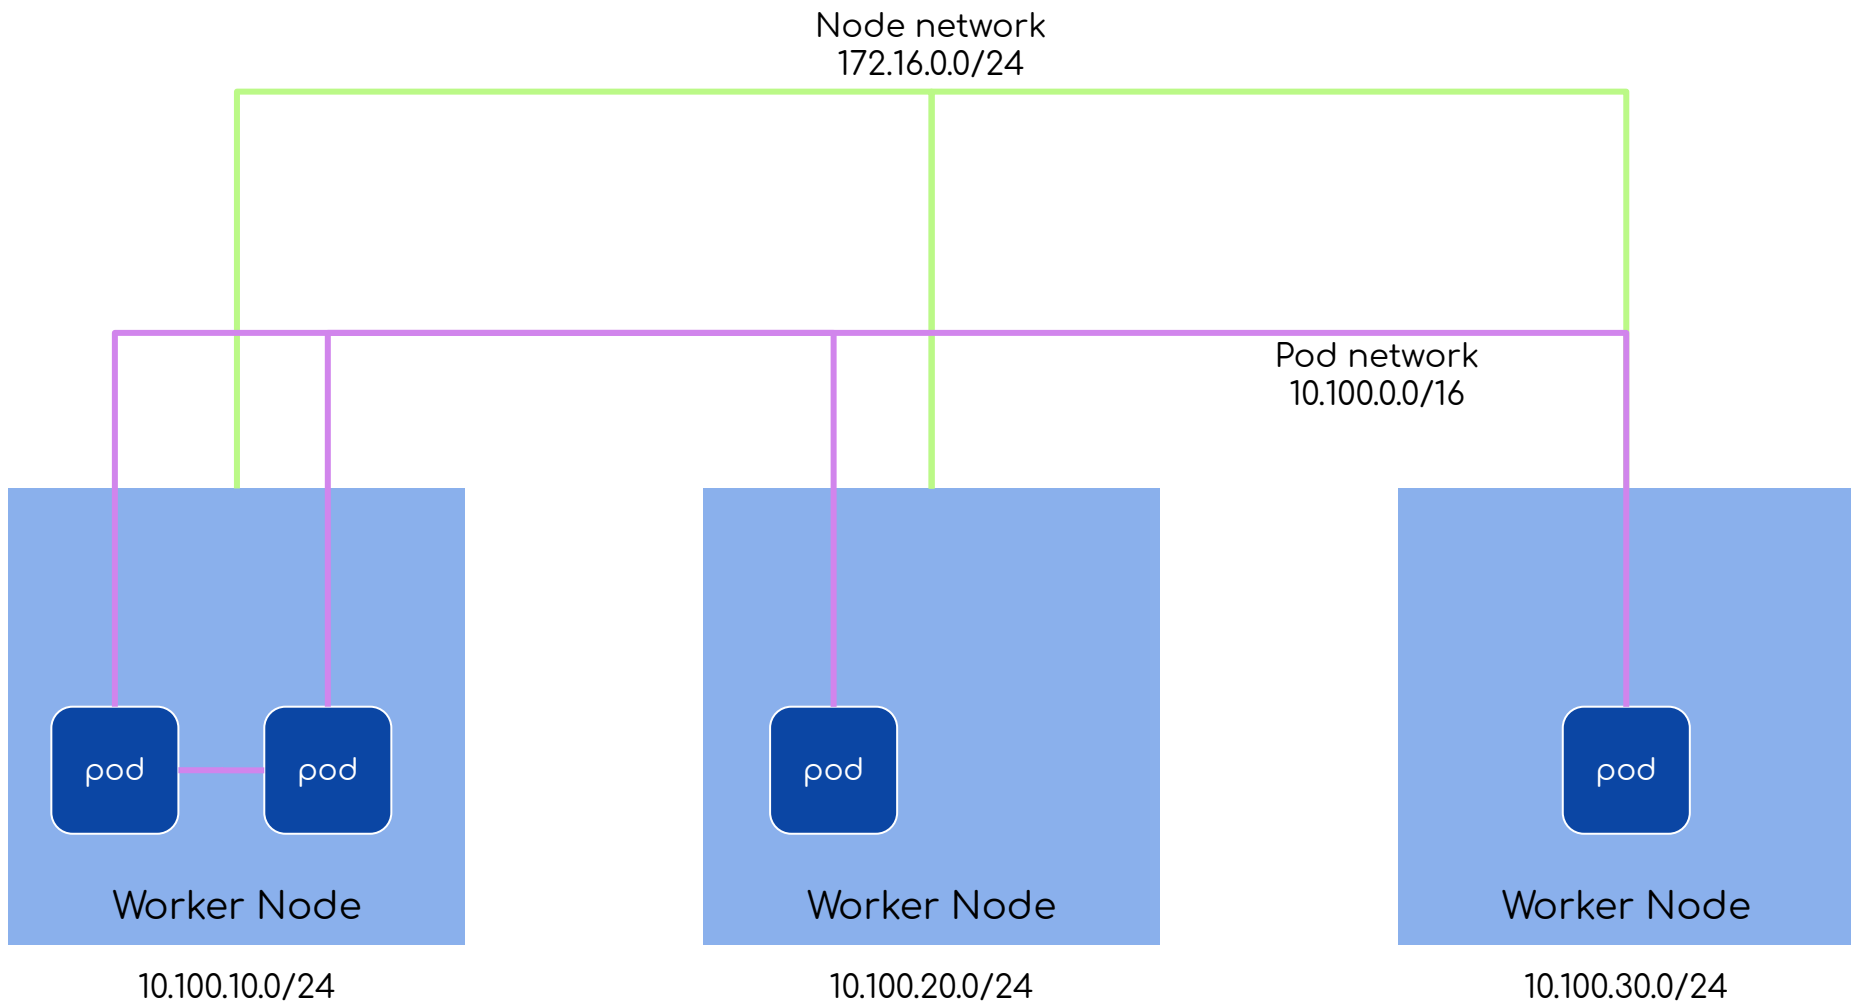
Node network  
172.16.0.0/24

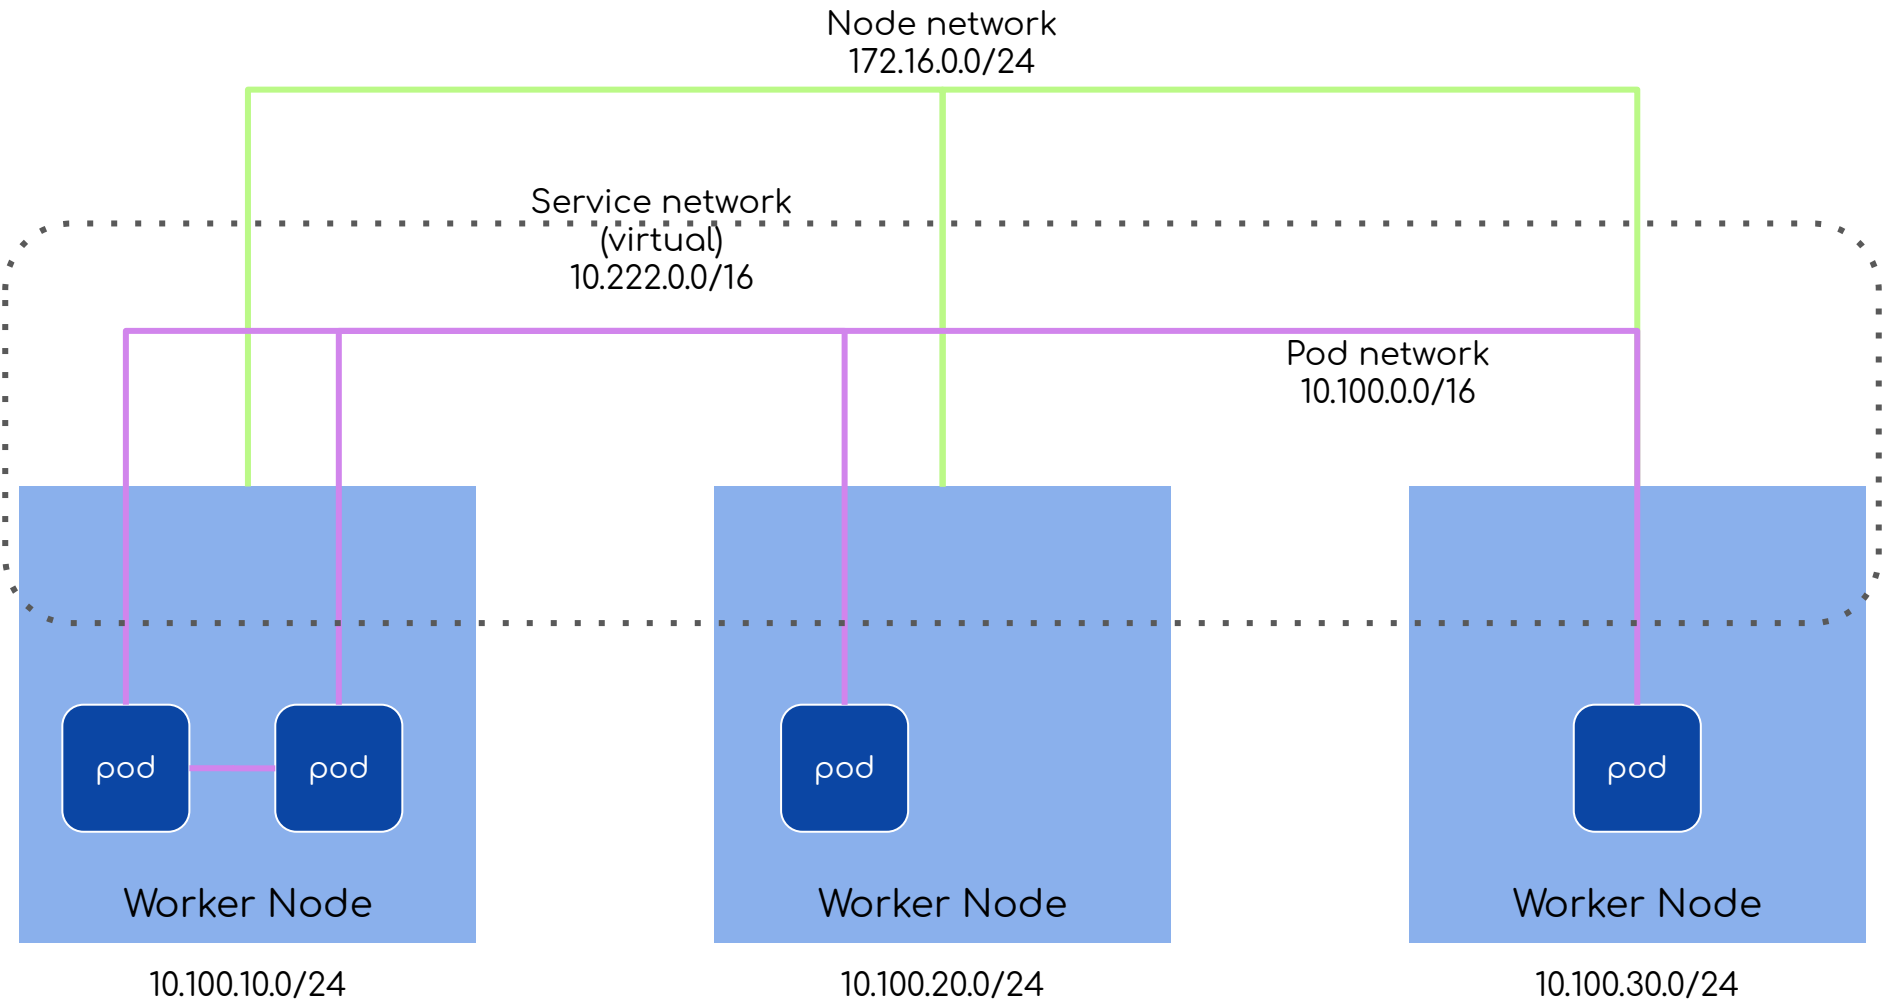



Worker Node

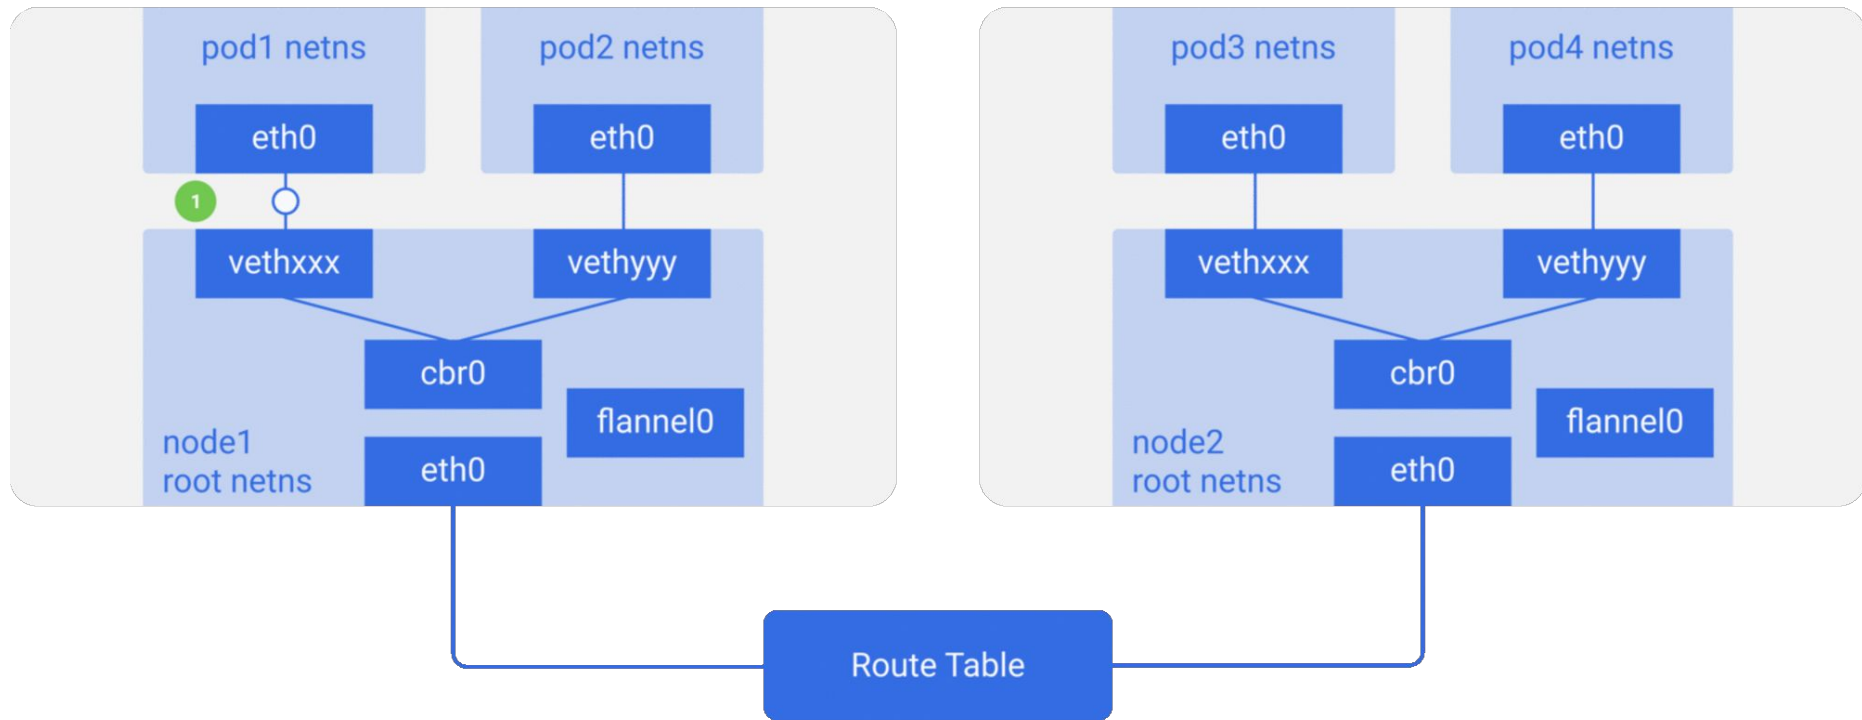
Worker Node

Worker Node





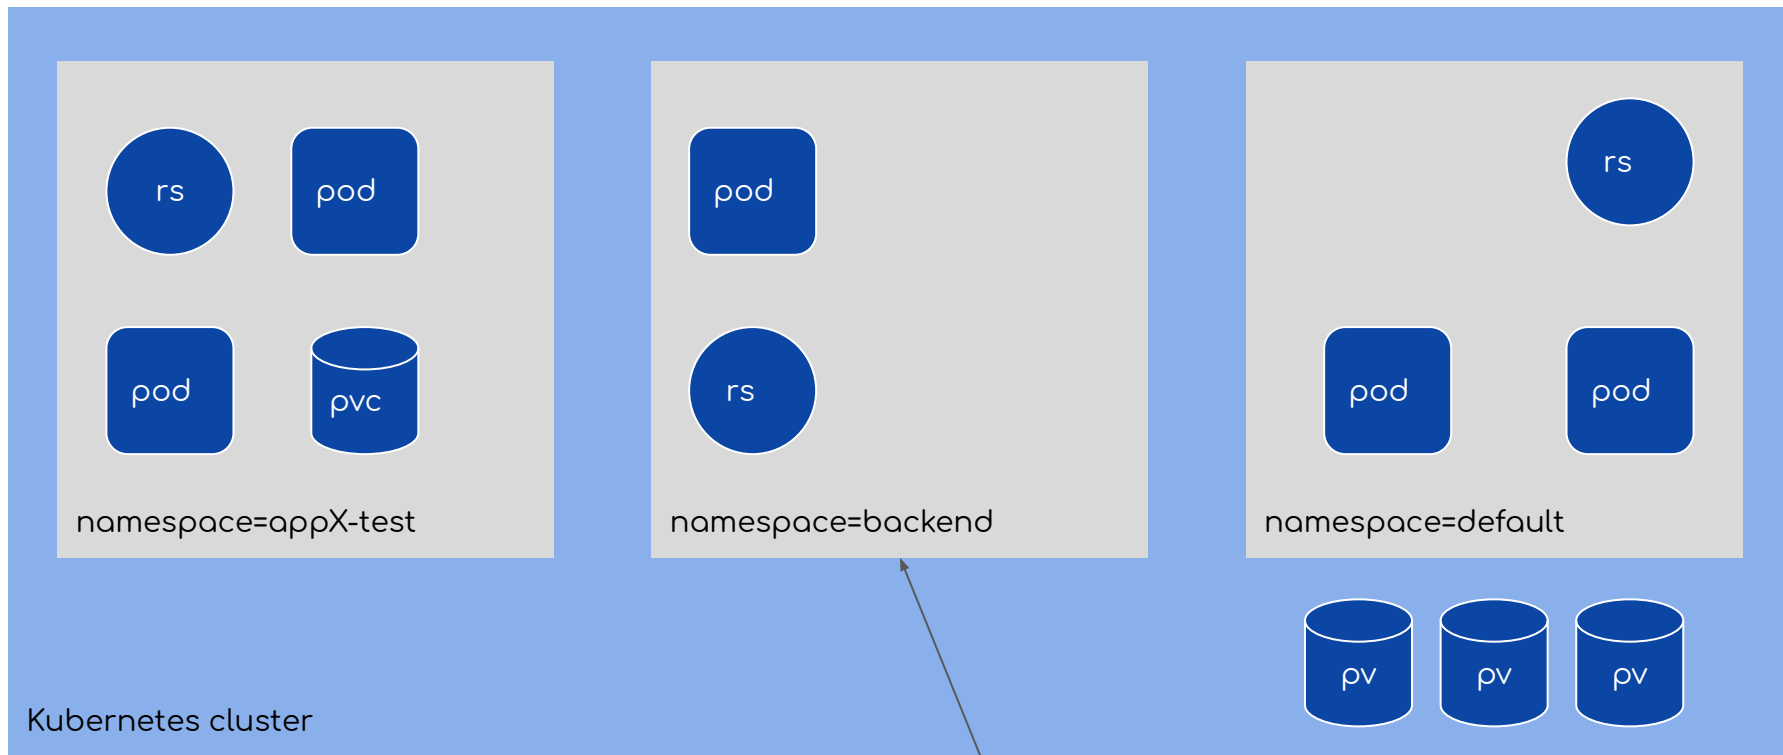




Source: <https://medium.com/@ApsOps/an-illustrated-guide-to-kubernetes-networking-part-2-13fdc6c4e24c>



# Kubernetes objects and namespaces



Kubernetes cluster

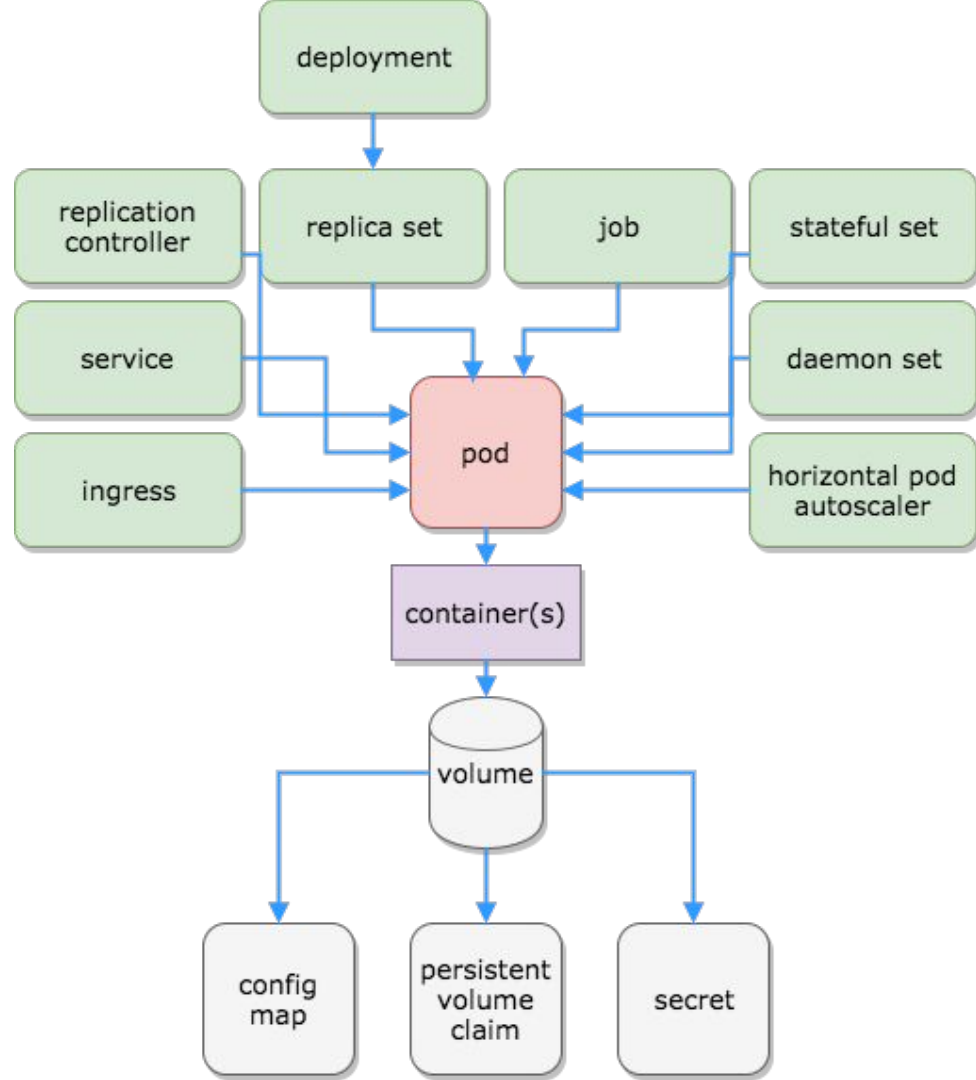
Namespaces



kubectl


~/kube/config:

- context:
  - cluster: ...
  - namespace: ...
  - user: ...




Kubernetes objects

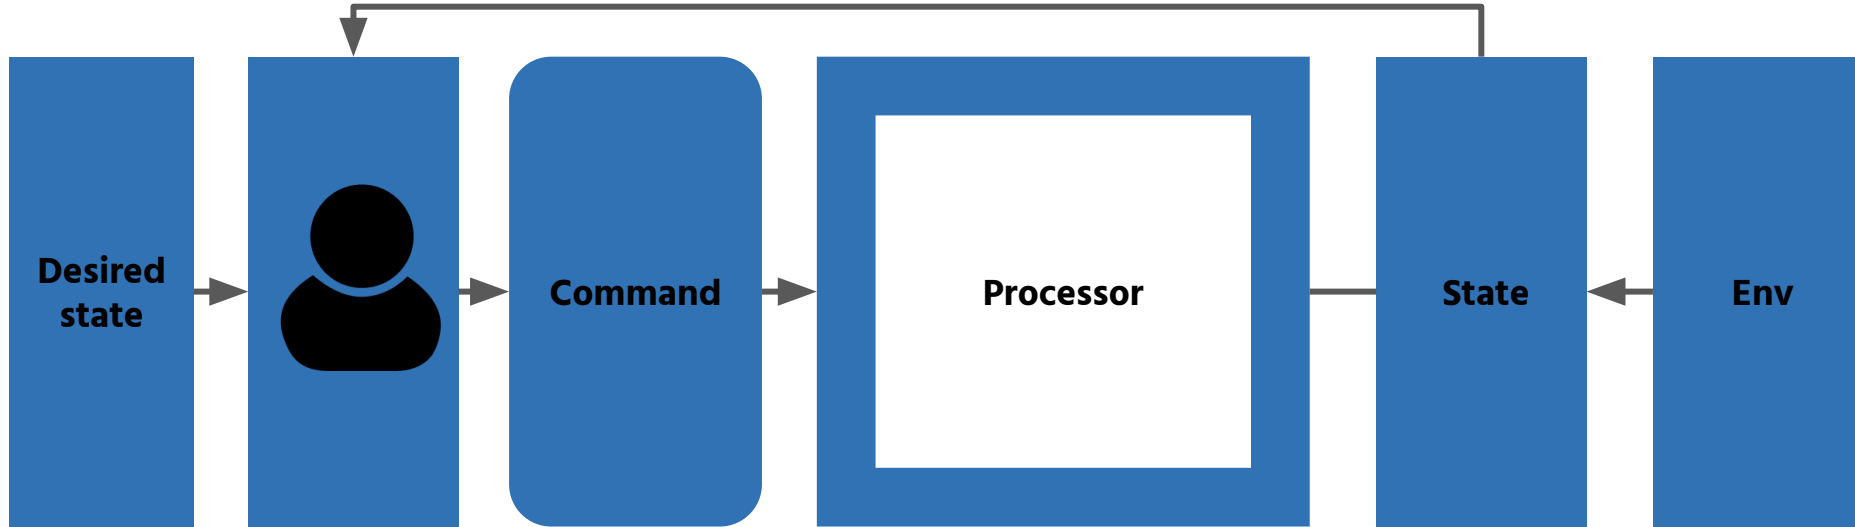
# Pod



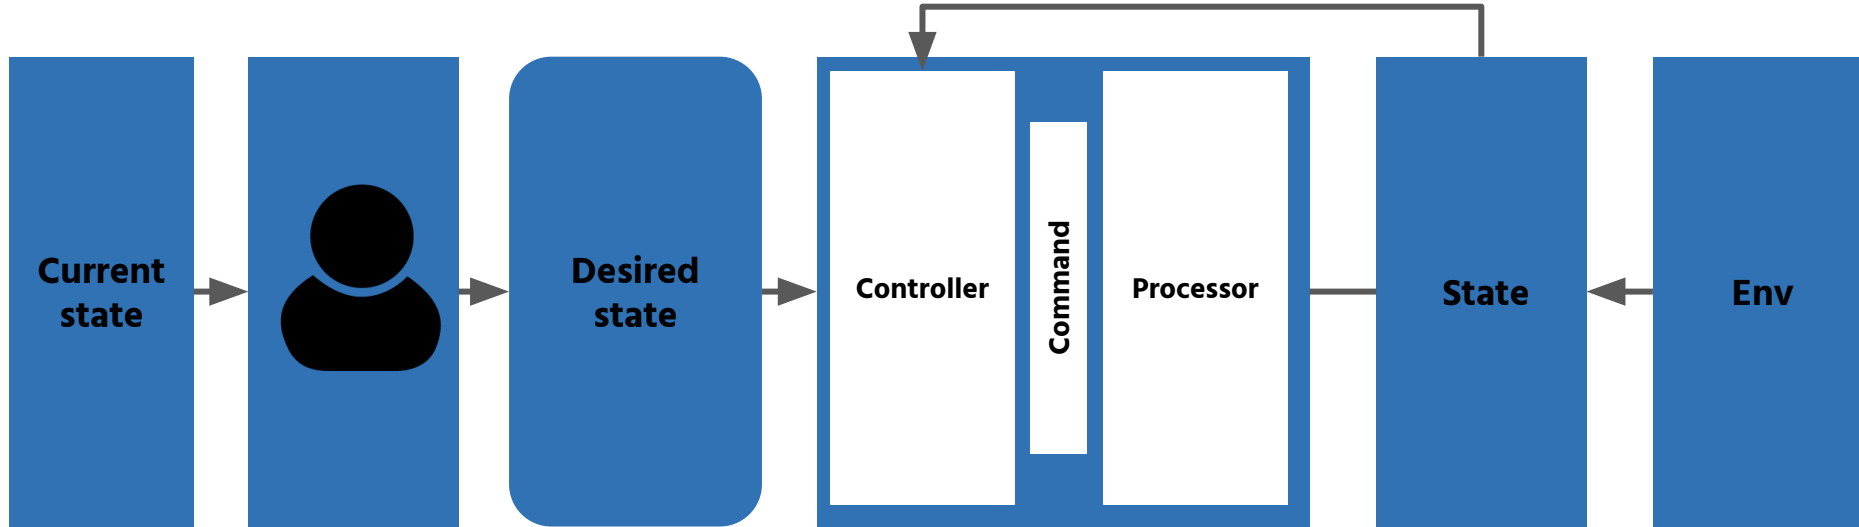
# Imperative vs. declarative



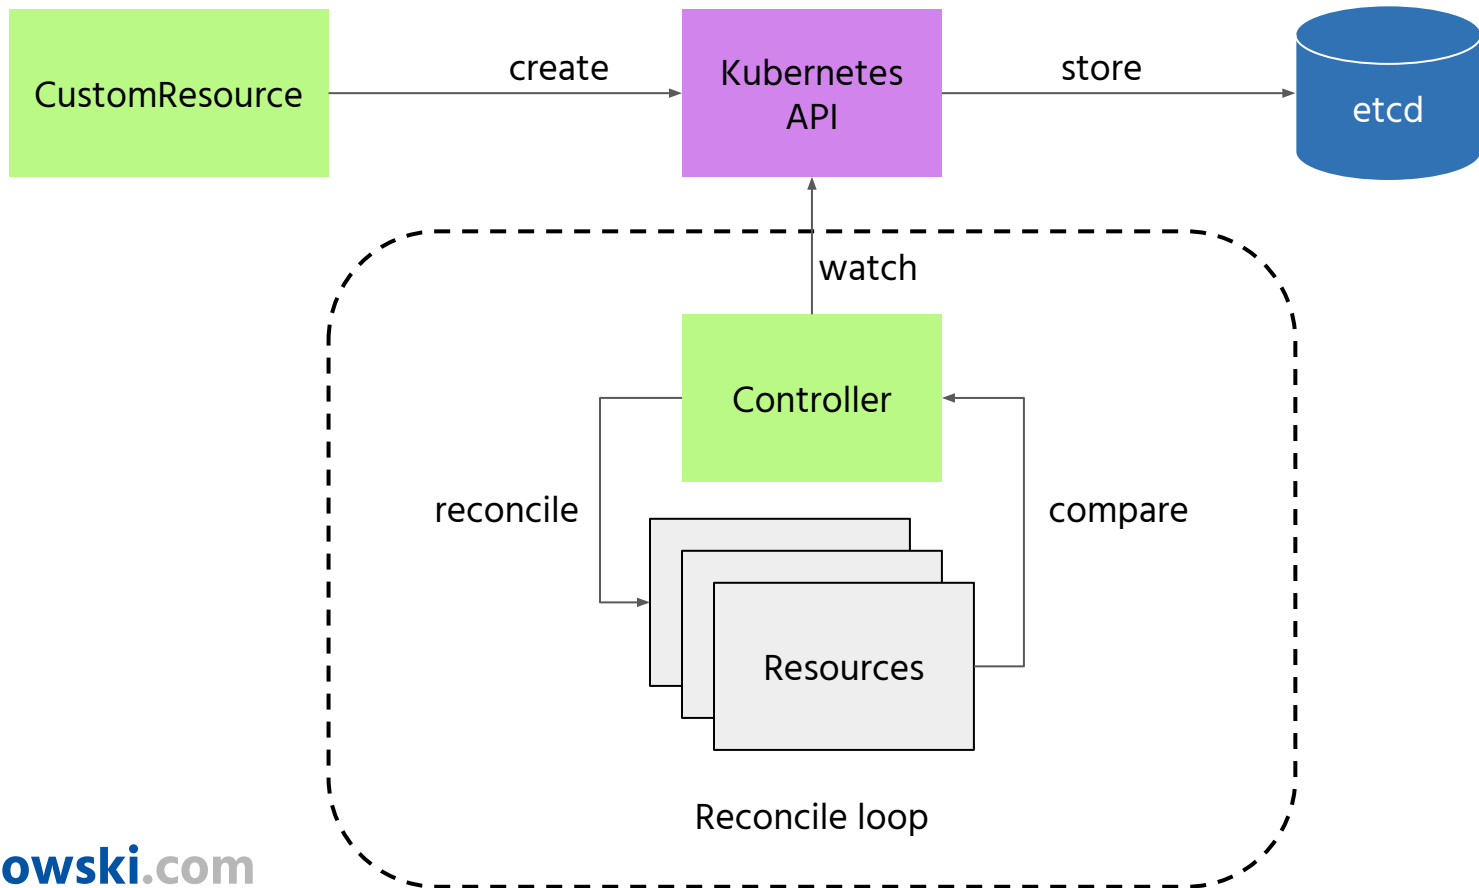
# Imperative systems



# Declarative systems



# Kubernetes controller with reconcile loop





# Yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: fussy
spec:
  containers:
  - image: nginx
    name: fussy
```

# Yaml

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: fussy
```

```
spec:
```

```
  containers:
```

```
  - image: nginx
```

```
    name: fussy
```

## REQUIRED

- different levels of stability and support
  - List with “*kubectl api-versions*”
  - alpha, beta, stable
- v1 = core group

# Yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: fussy
spec:
  containers:
  - image: nginx
    name: fussy
```

REQUIRED

# Yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: fussy
spec:
  containers:
    - image: nginx
      name: fussy
```

REQUIRED

Valuees required: **name**

Optional: namespace

# Yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: fussy
  labels:
    app: fussy
spec:
  containers:
```

OPTIONAL

Key-Value pairs

# Yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: fussy
spec:
  containers:
  - image: nginx
    name: fussy
```

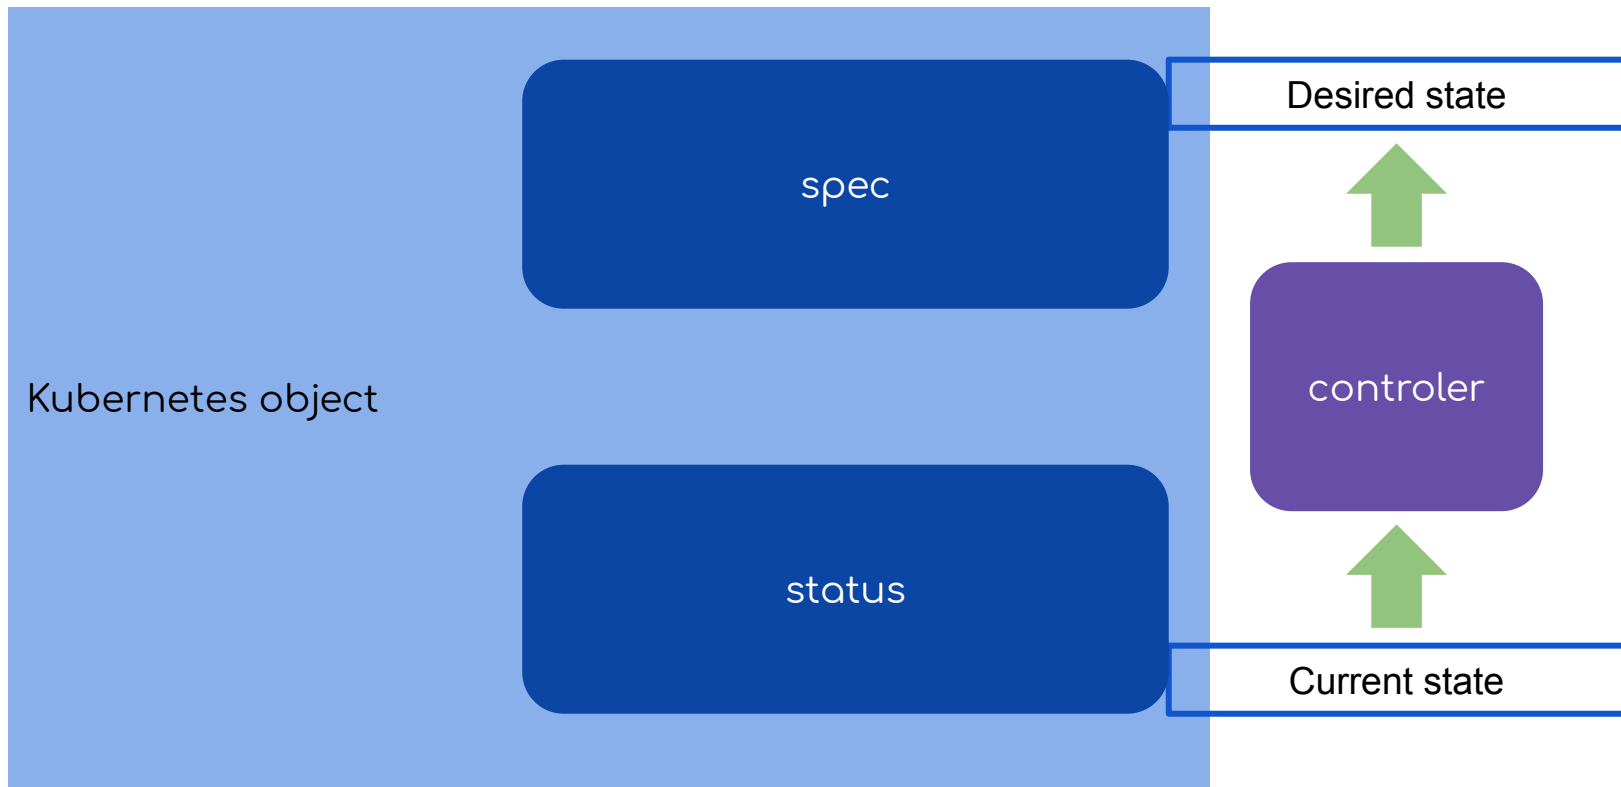
OPTIONAL

# Yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: fussy
spec:
  containers:
  - image: nginx
    name: fussy
```

OPTIONAL

Depends on object kind



Declarative nature of objects



# Labels vs. Annotations

## Labels

- Used to organize and select subsets of objects
- Key: Value
- Used internally and by clients


## Annotations

- Used for non-identifying, metadata objects
- Structured, unstructured data
- Used by libraries, clients

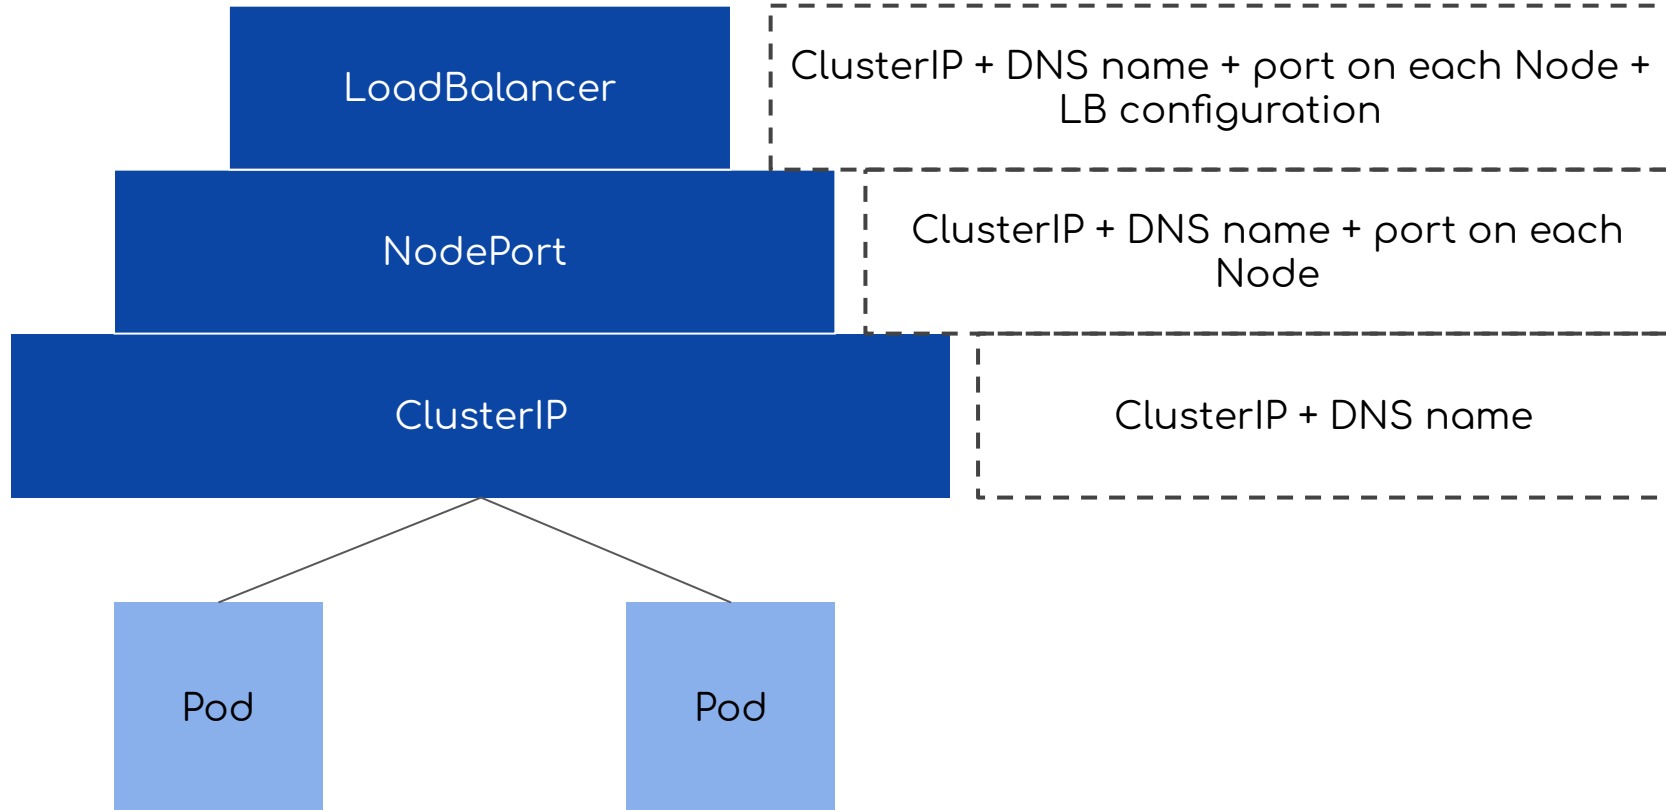
# Service

## Types

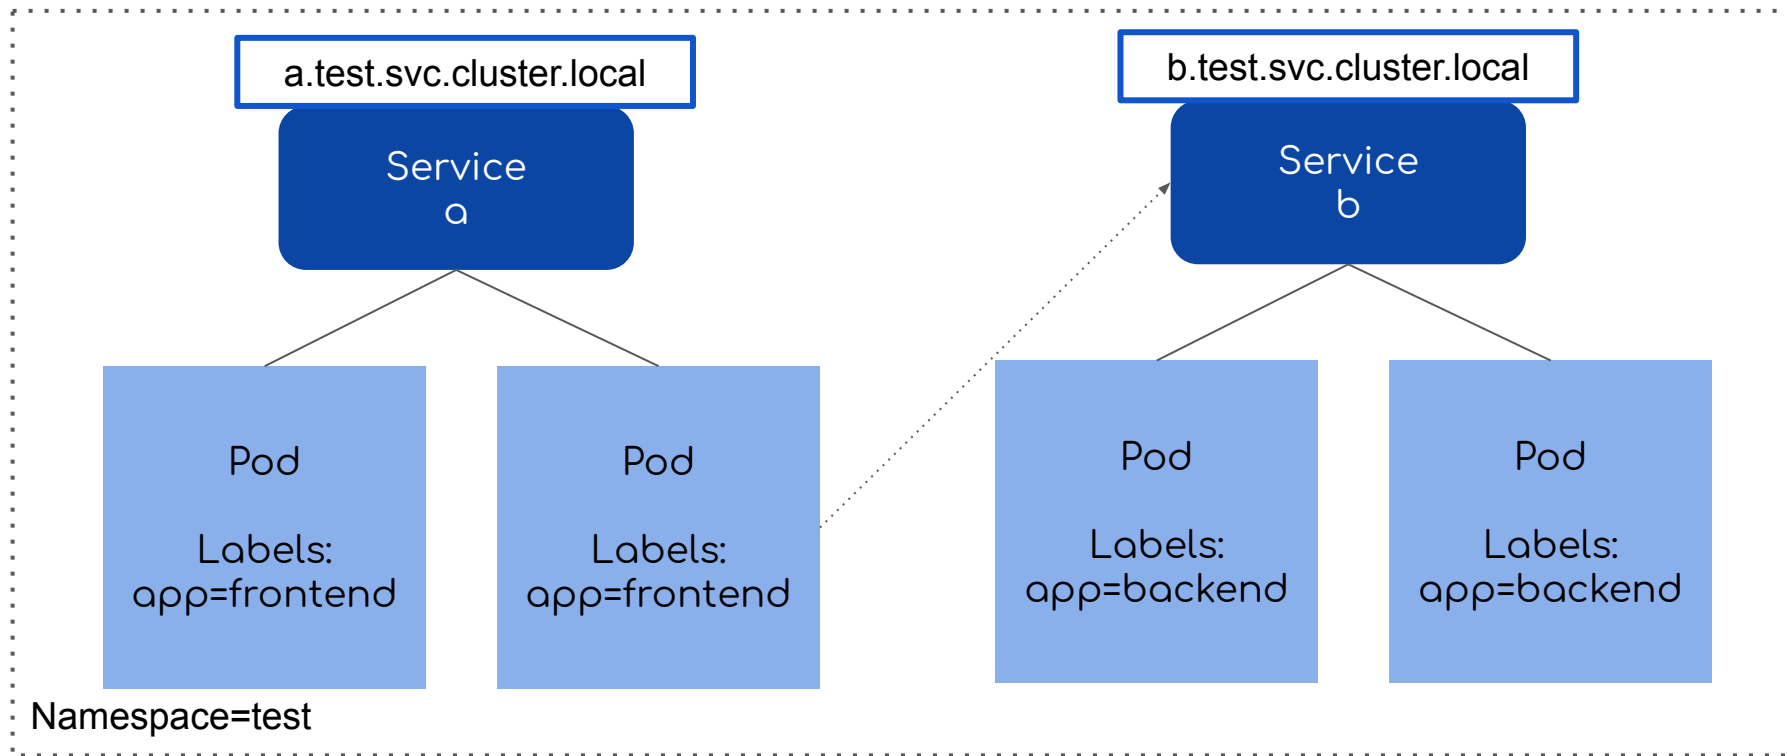
- ❑ ClusterIP
- ❑ NodePort
- ❑ LoadBalancer
- ❑ ExternalName



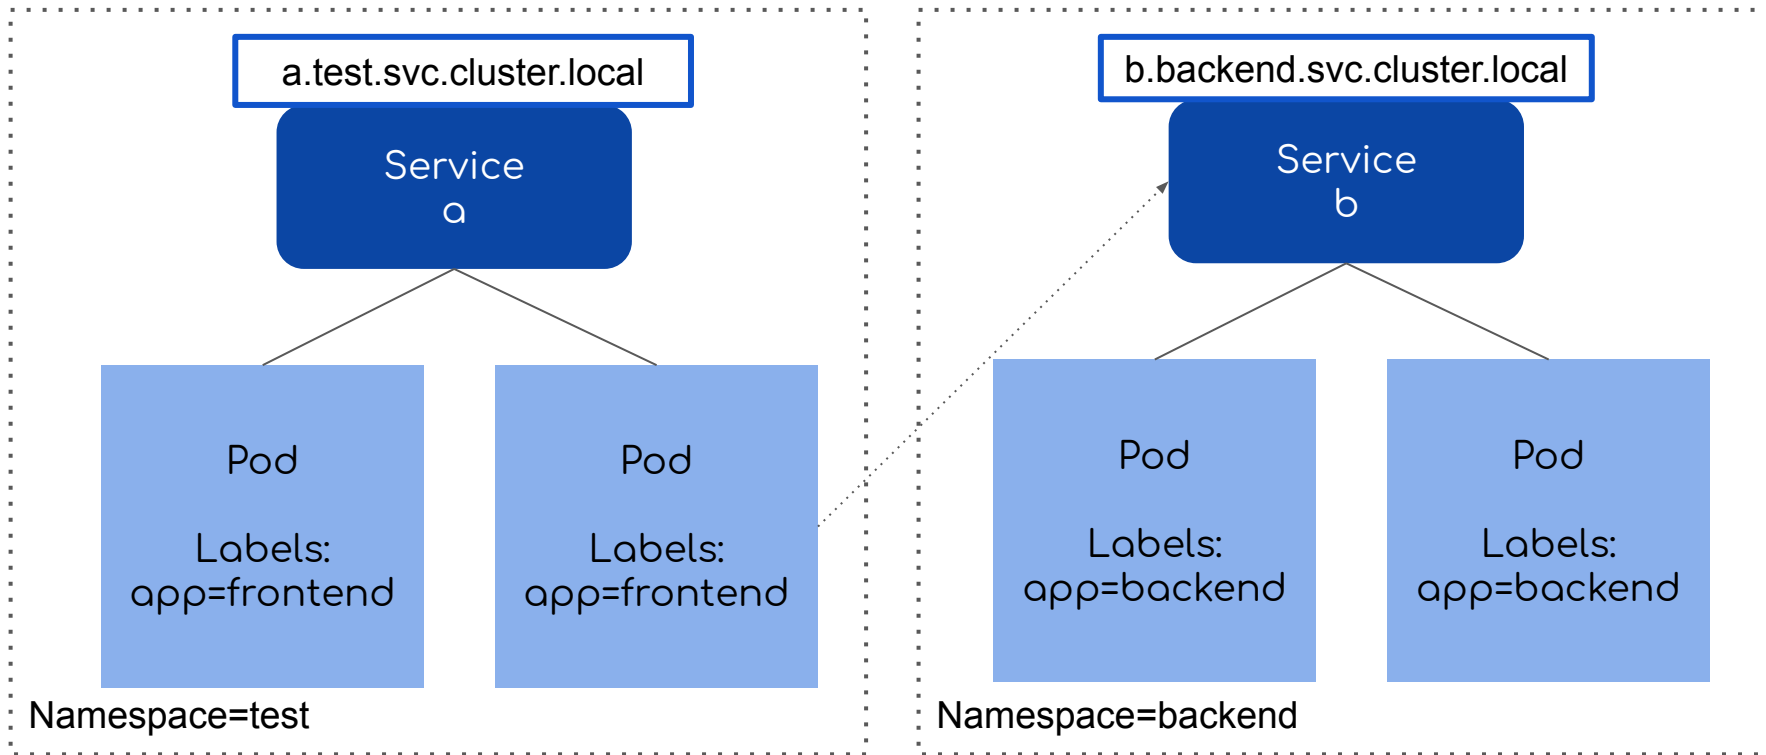
# Service types



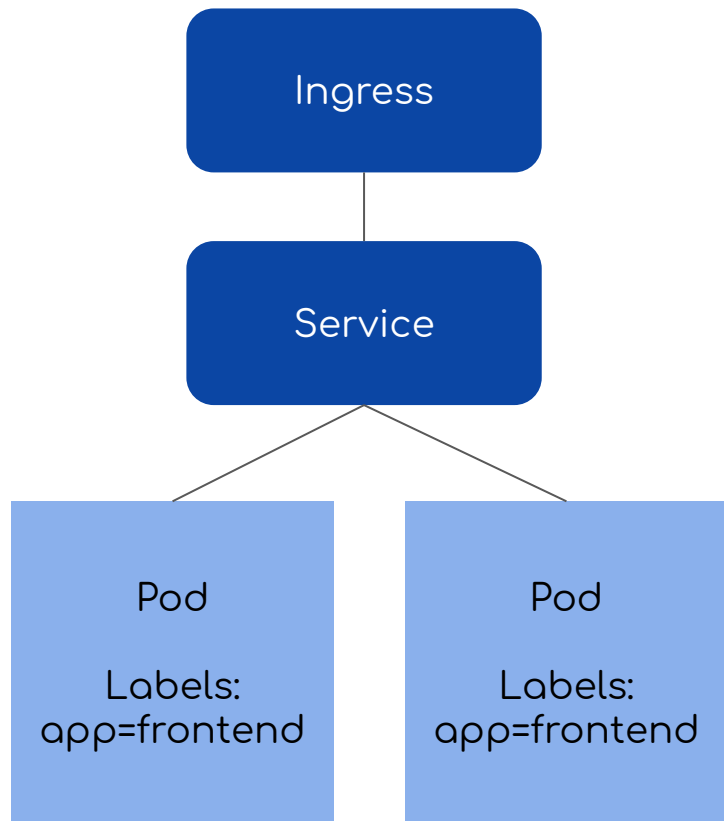
# Service Discovery



# Service Discovery



# Ingress




Interface for http(s) traffic


Best known implementation based on Nginx

Controlled fully with declarative code and annotations

# ConfigMap - as environment variables




# ConfigMap - as volumes






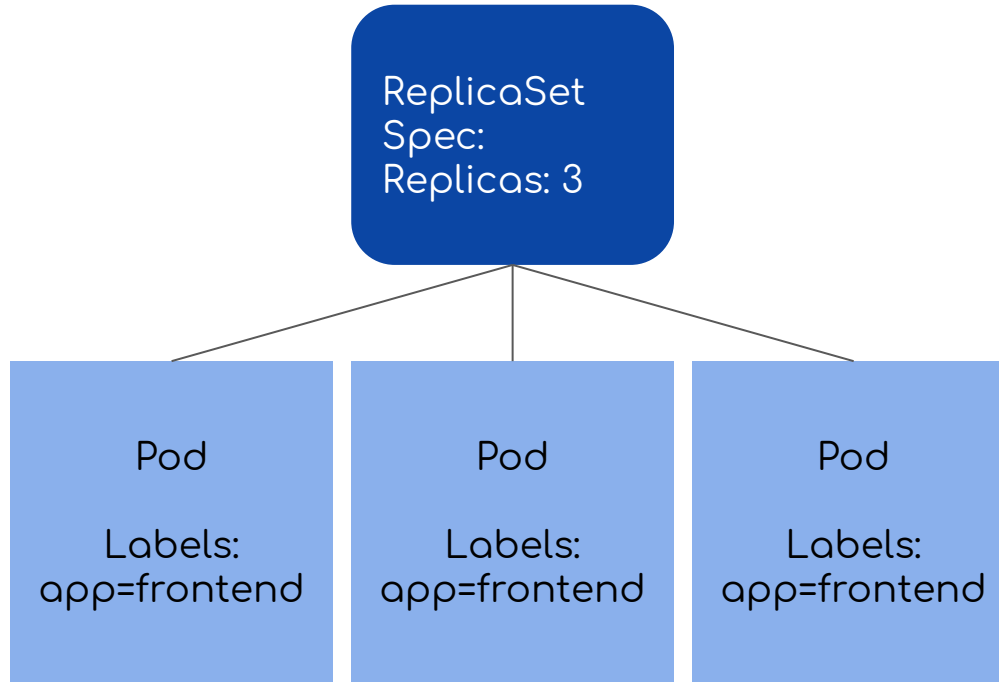
# ConfigMap - as volumes



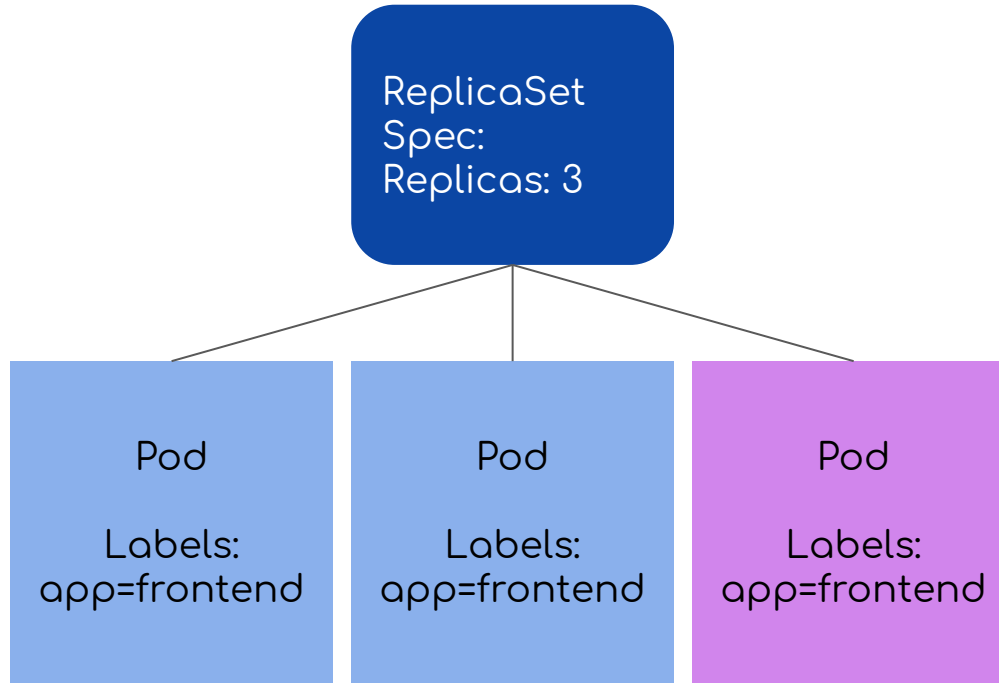
# Secrets



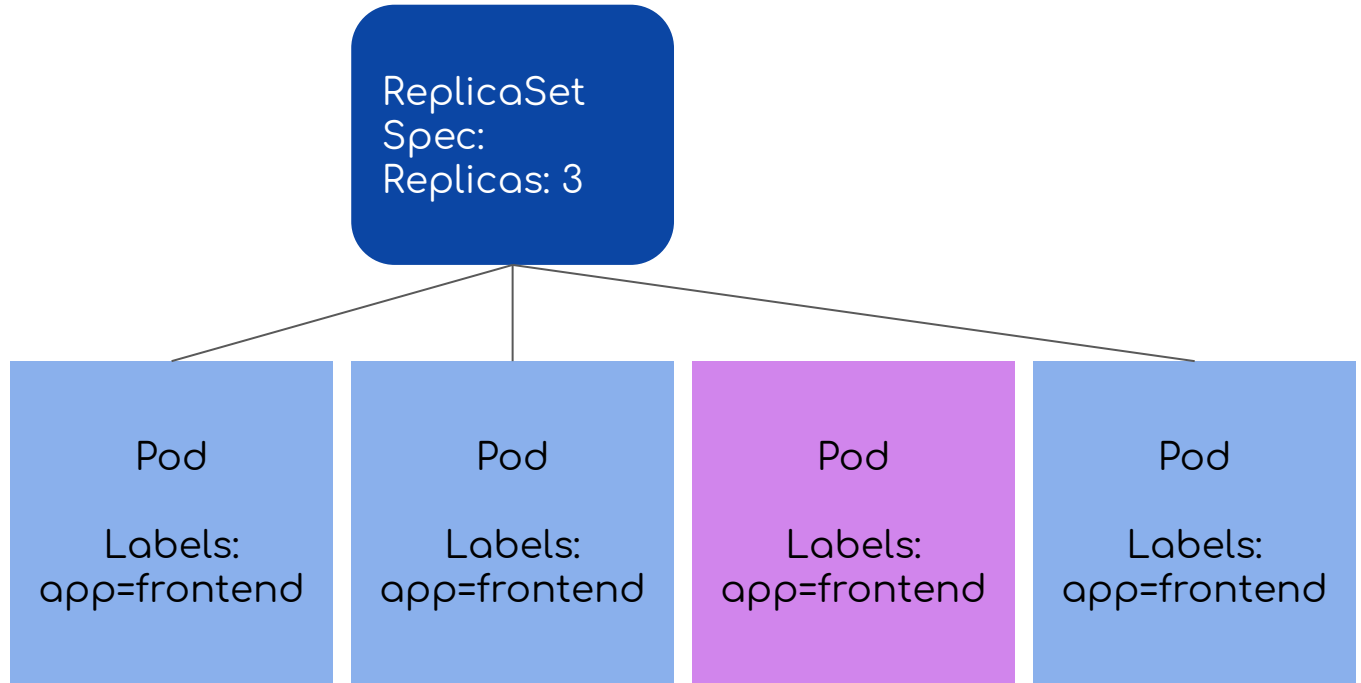
# ReplicaSet



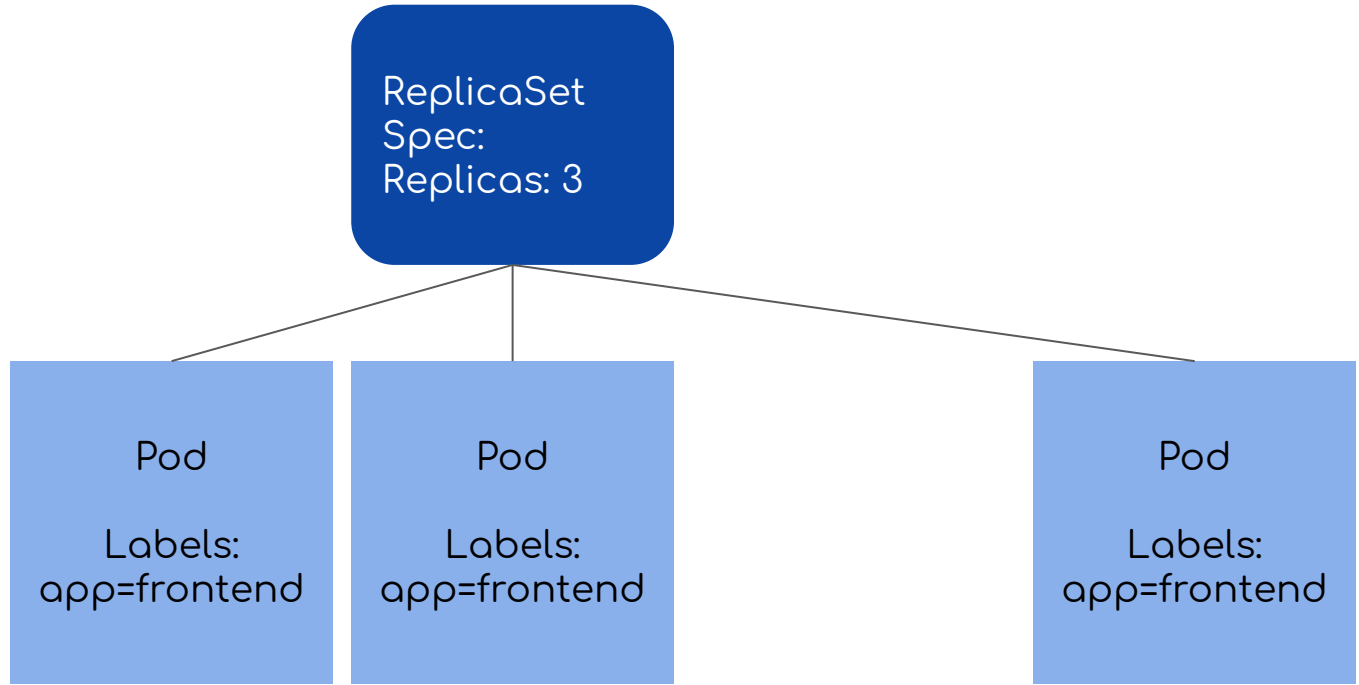
# ReplicaSet




# ReplicaSet



# ReplicaSet



# Volumes



Persistent  
Volumes

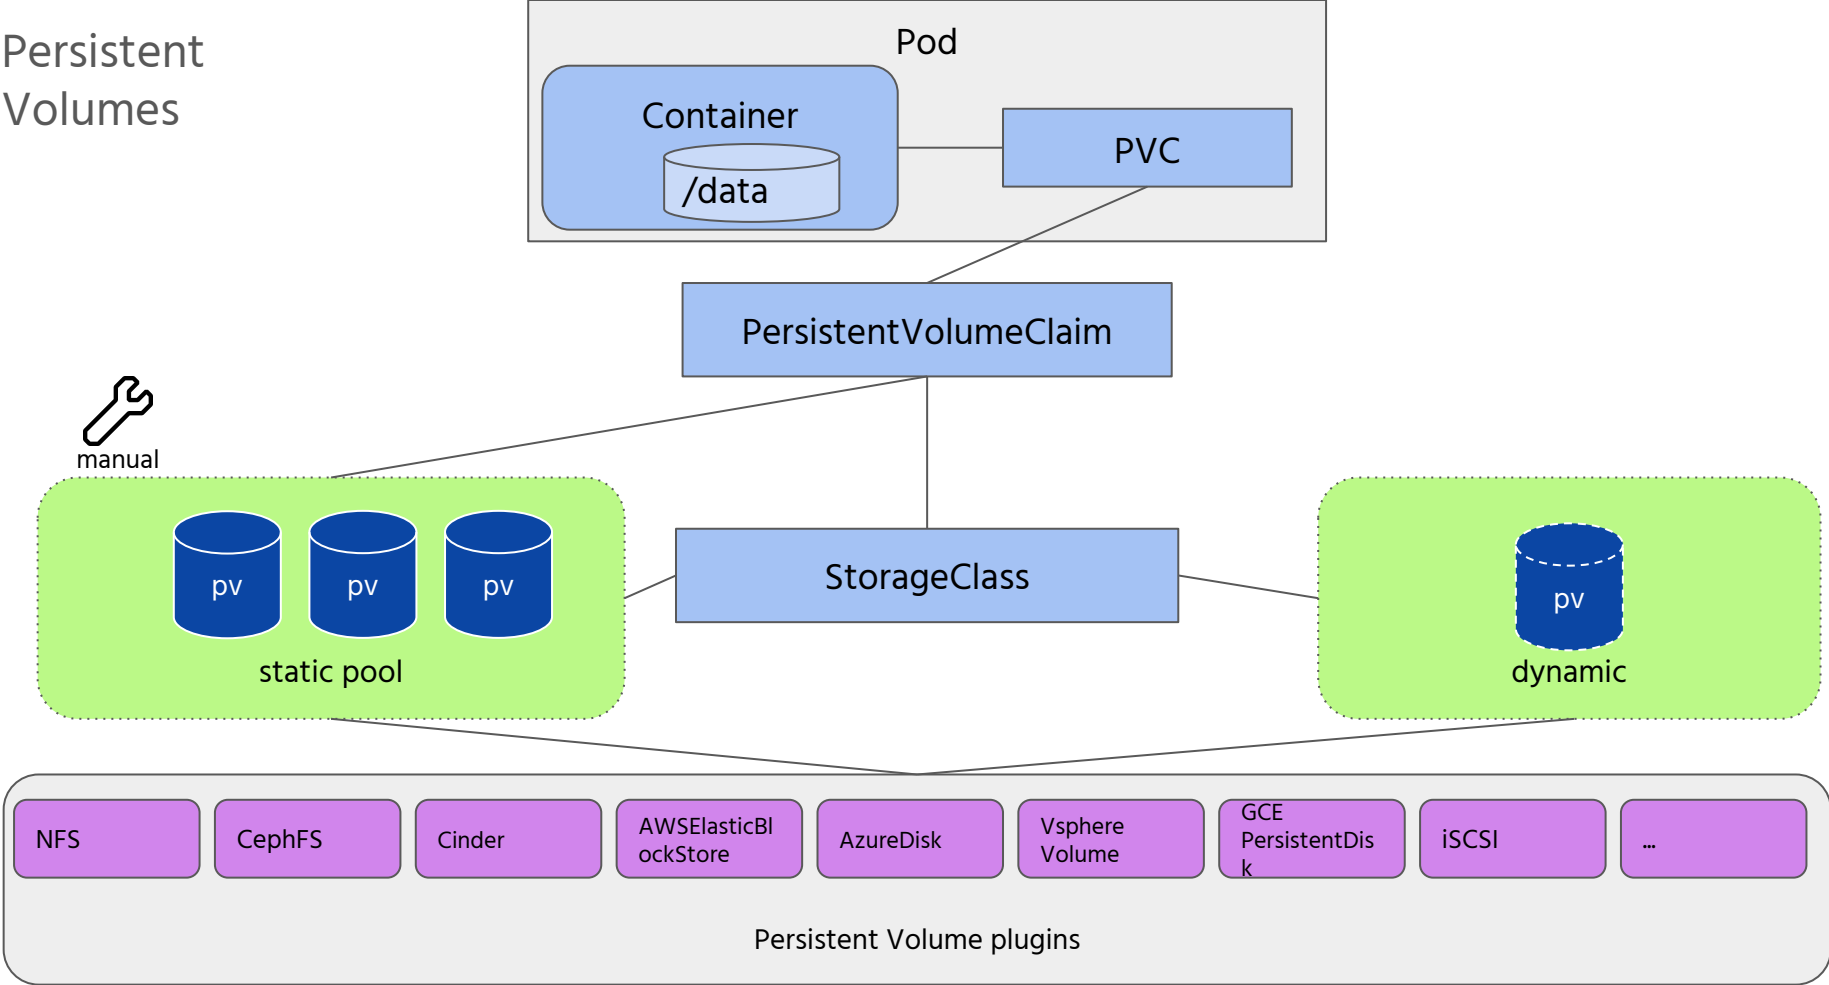
*Where*

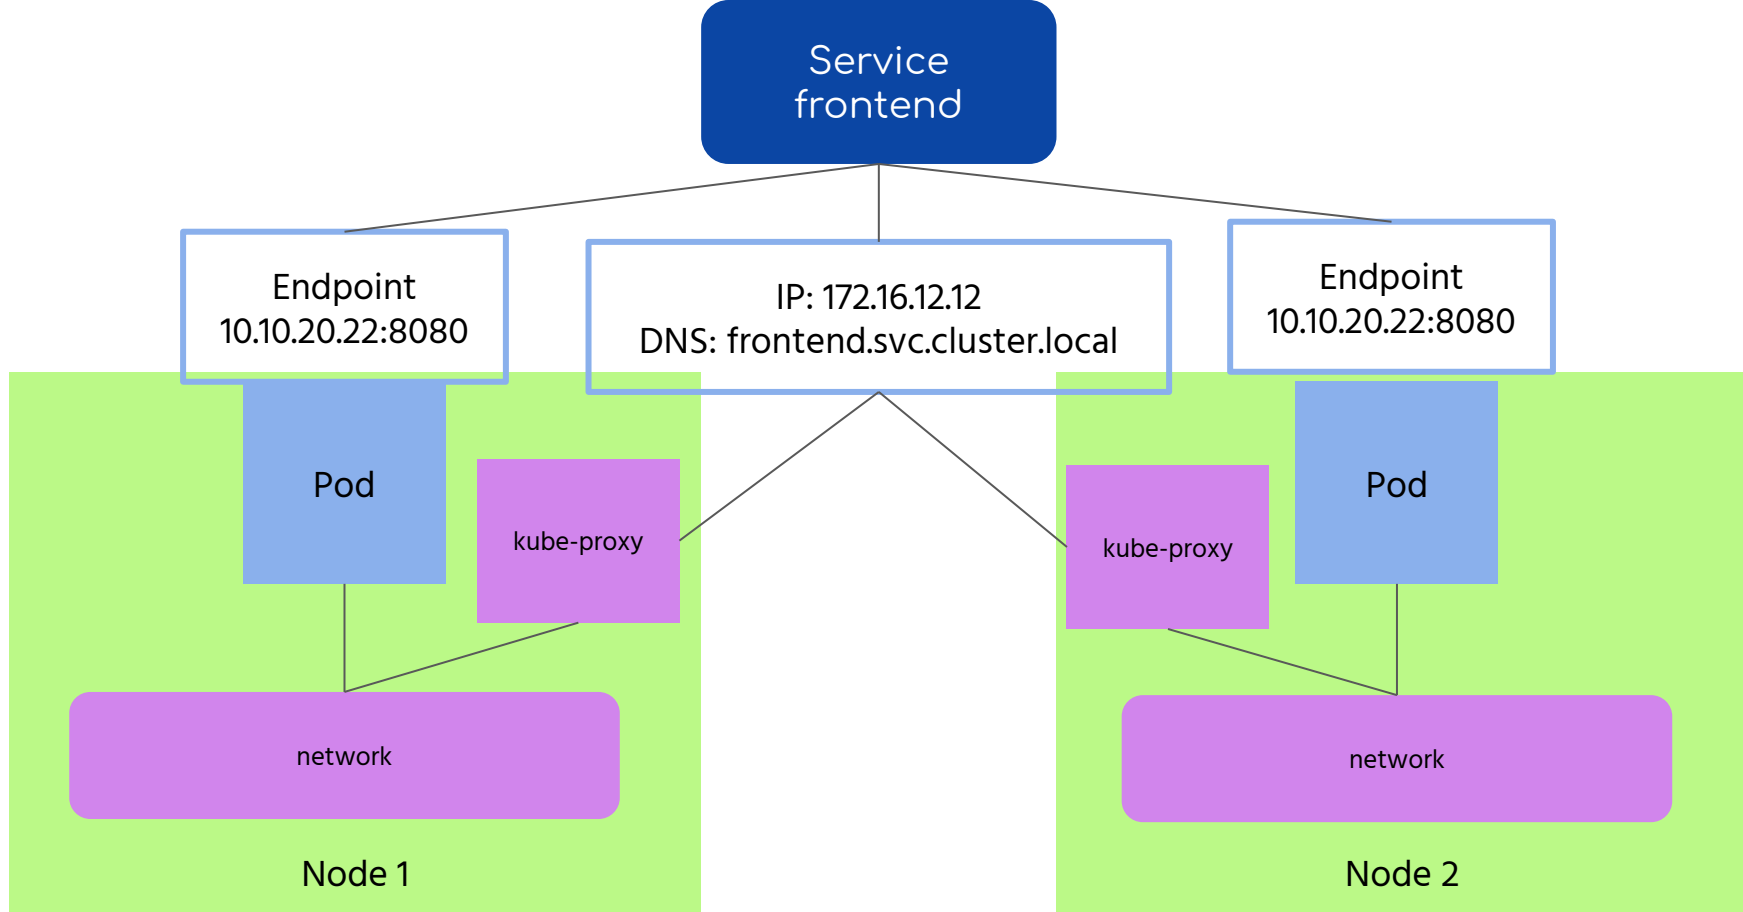
*What*

*How*



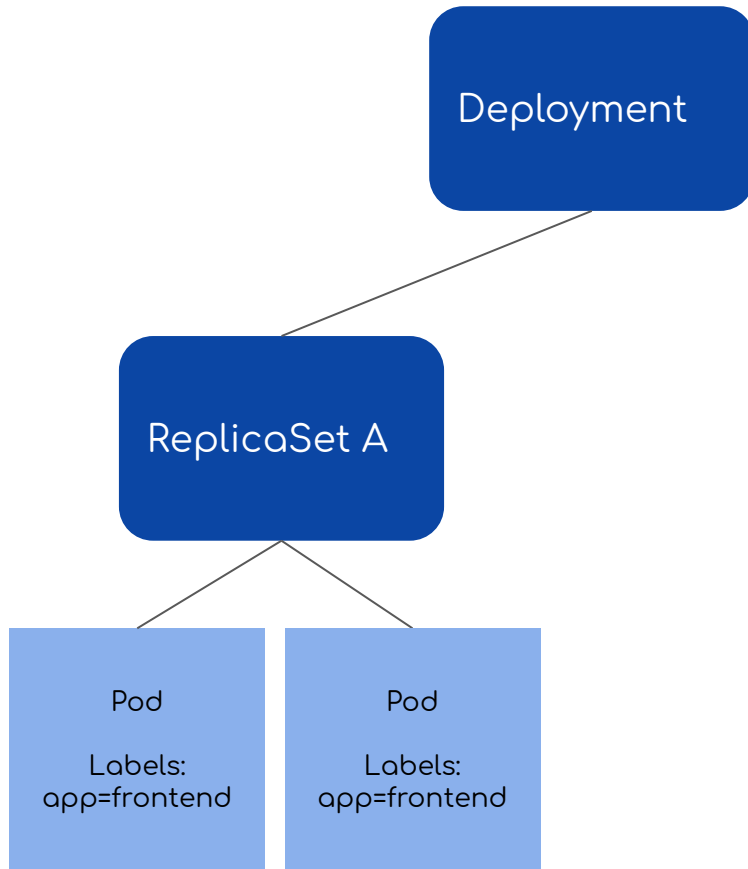
# Persistent Volumes



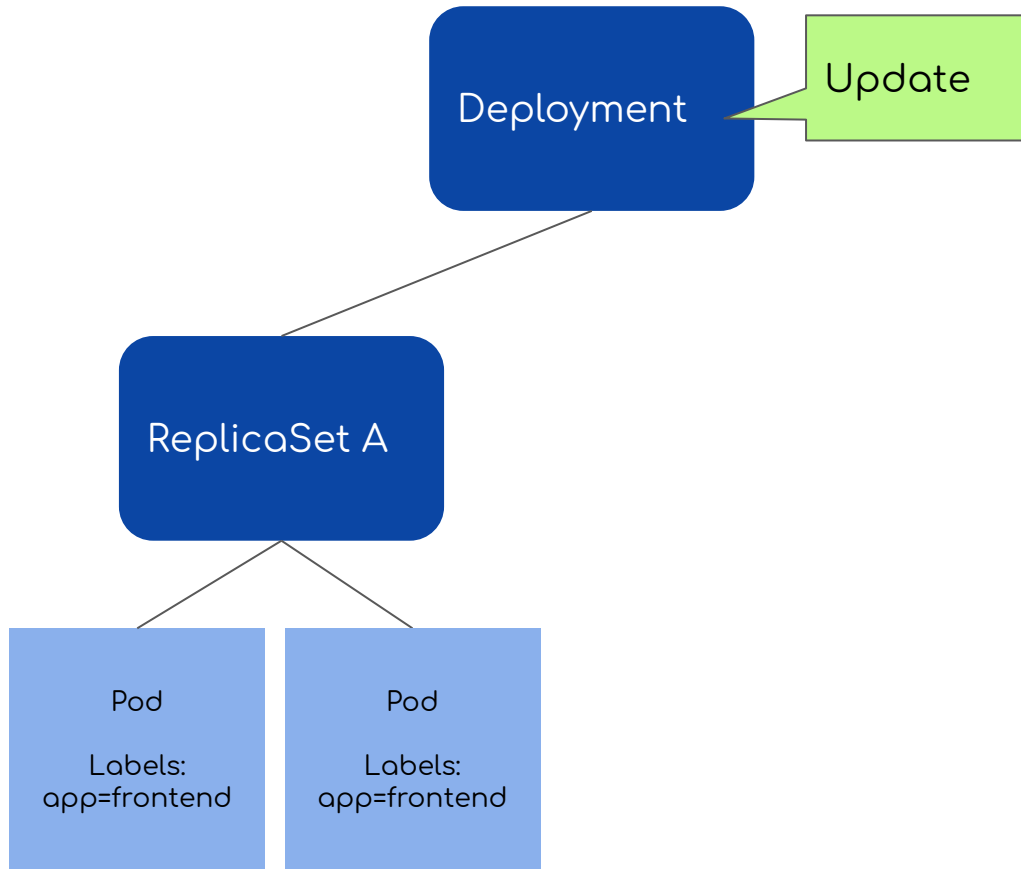


Service implementation details

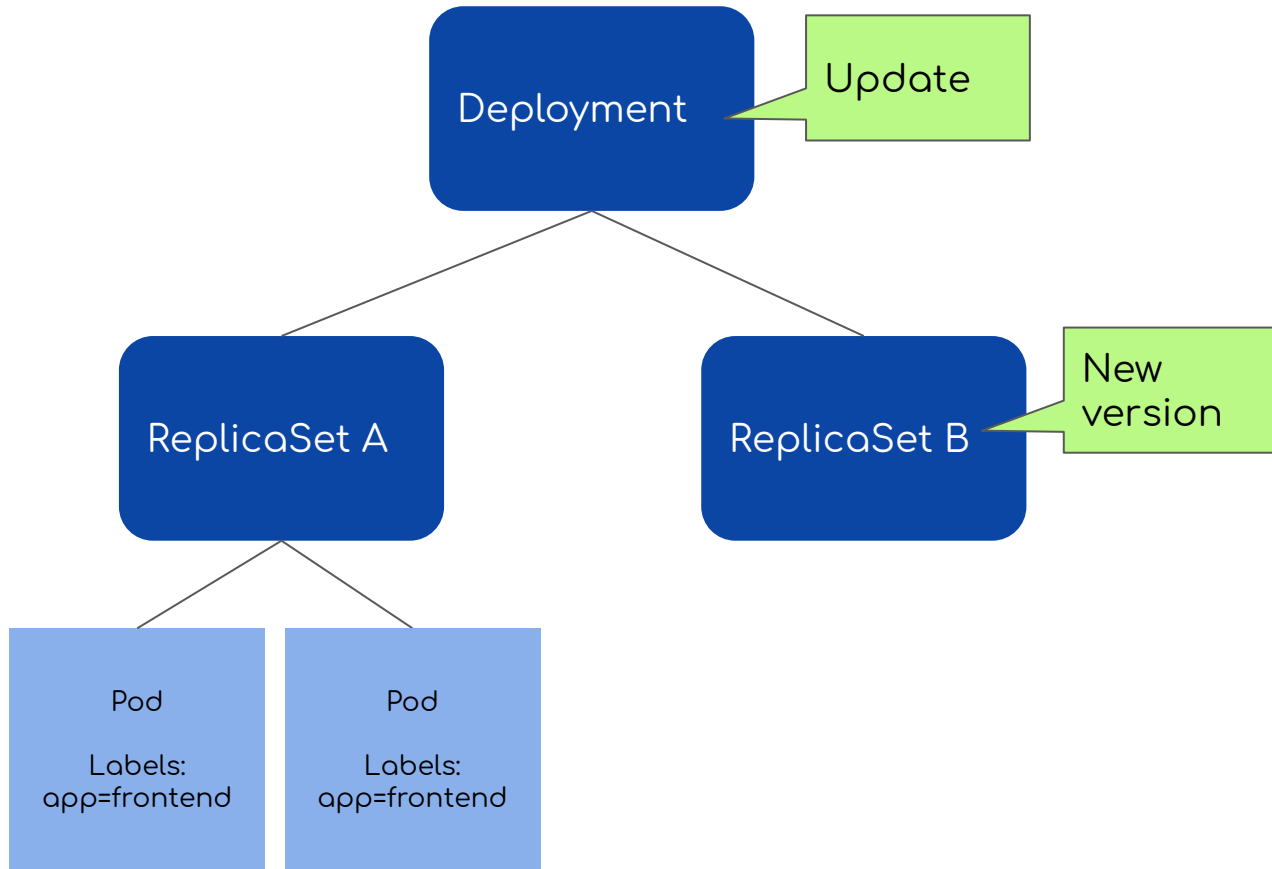
# Deployment with rolling update



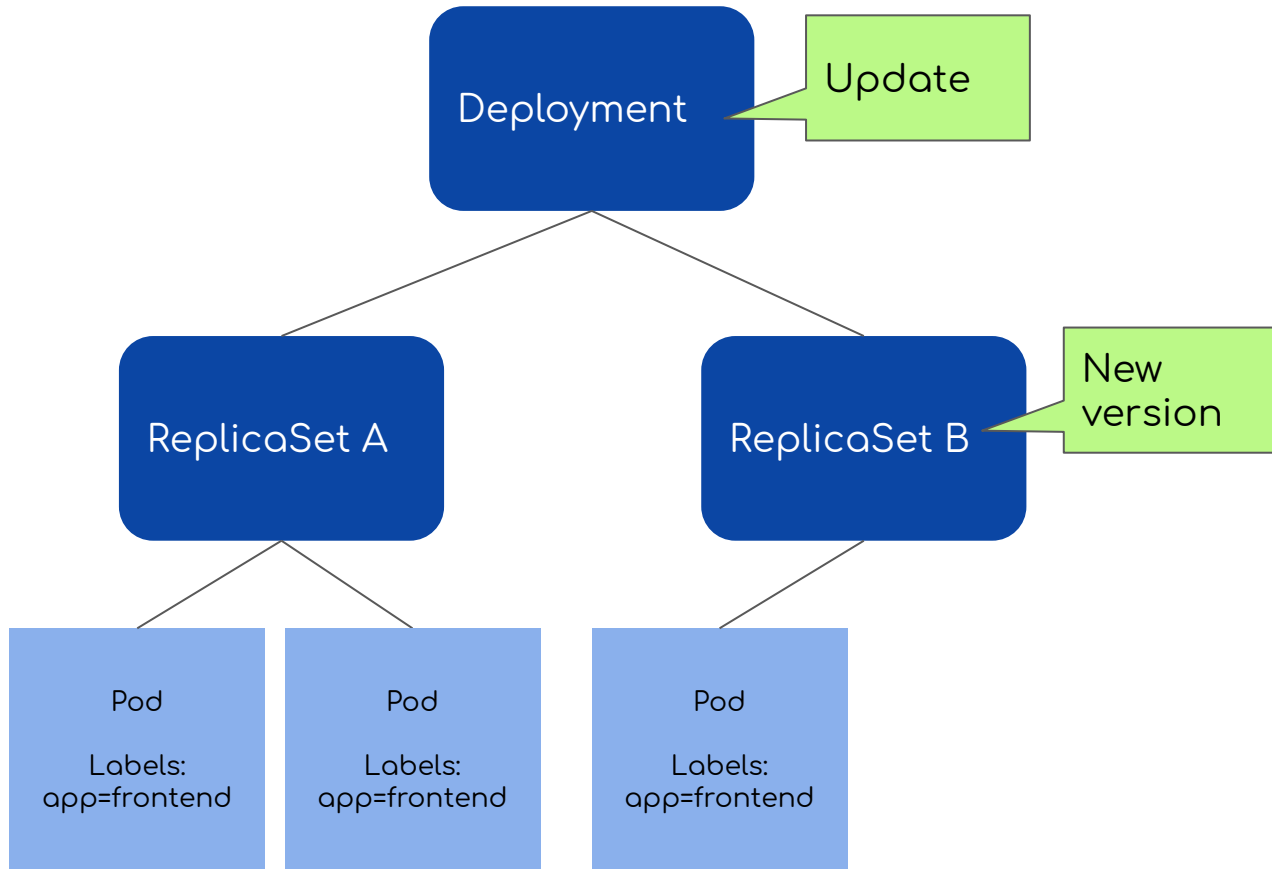
# Deployment with rolling update



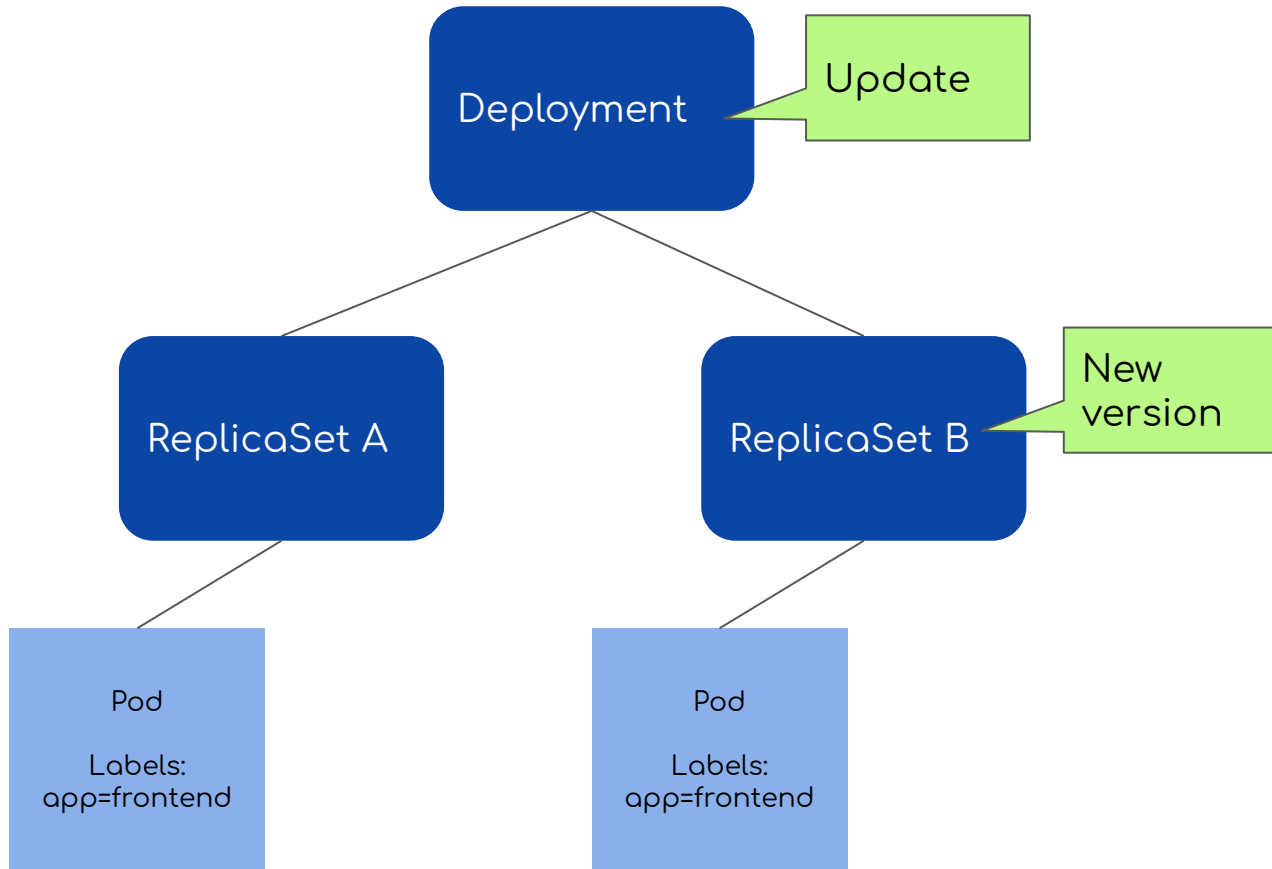
# Deployment with rolling update



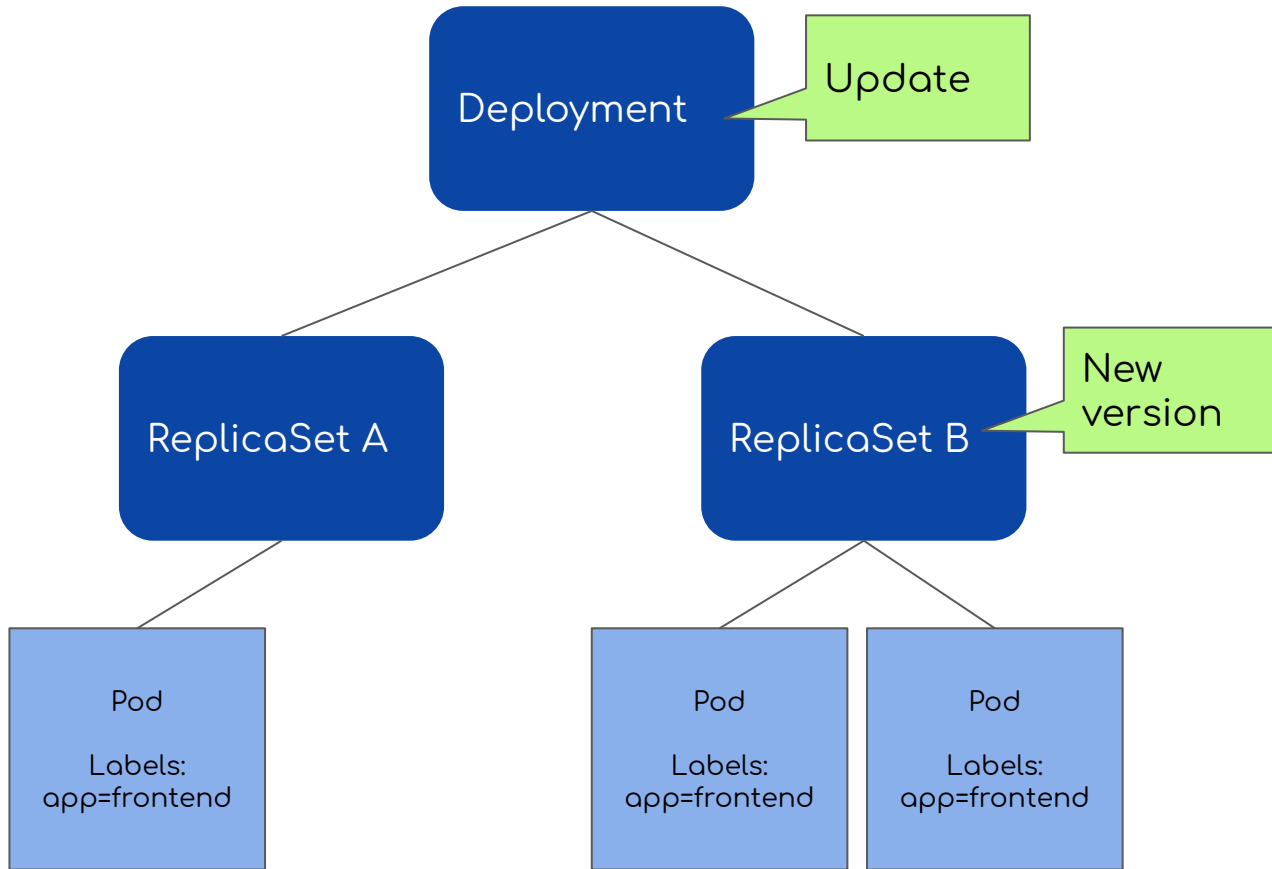
# Deployment with rolling update



# Deployment with rolling update

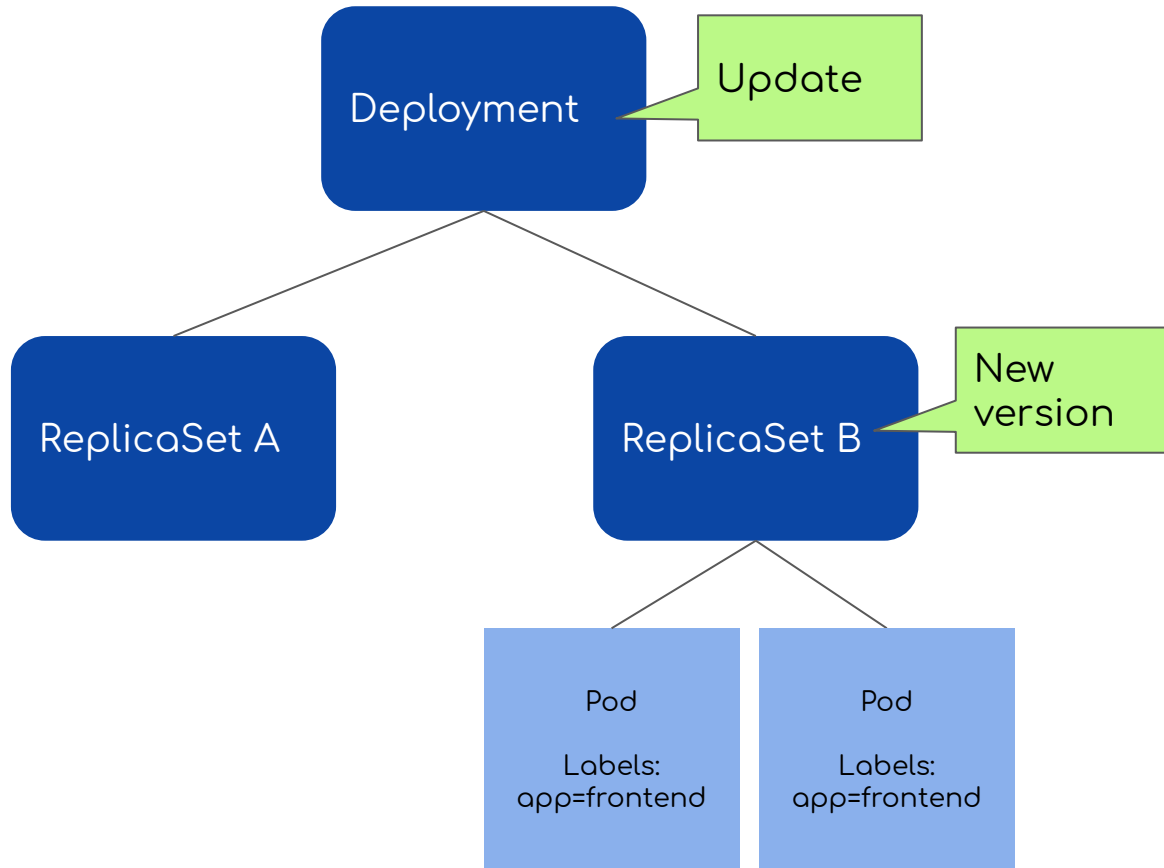


# Deployment with rolling update







# Deployment with rolling update




# Deployment strategies - Rolling Deployment




# Deployment strategies - Recreate

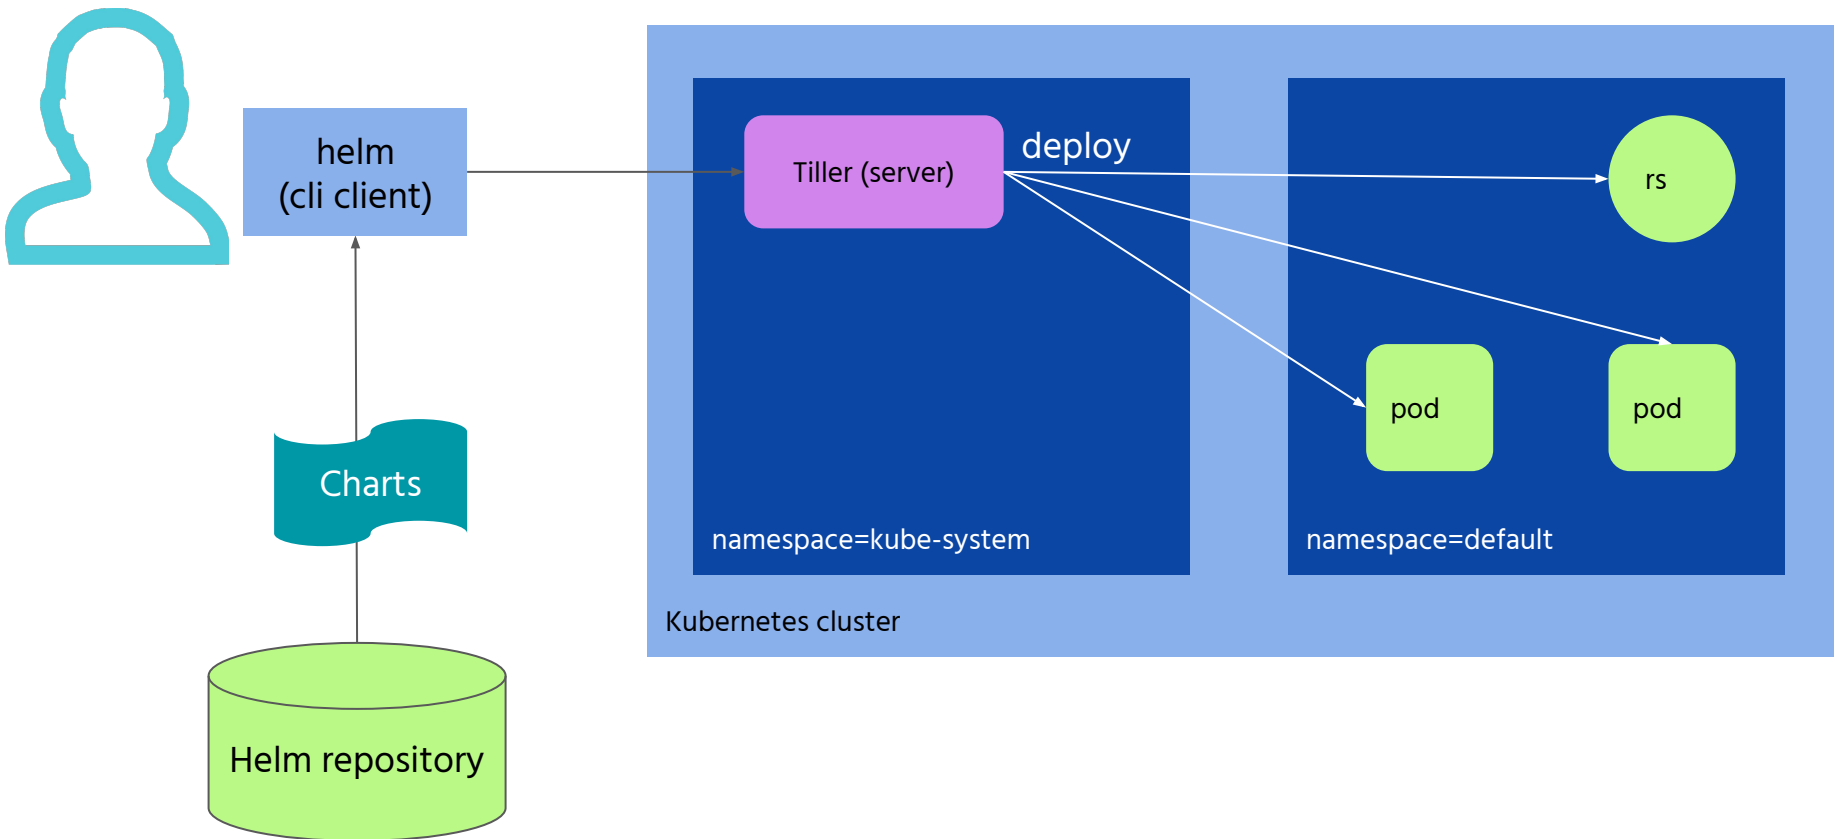


# Deployment strategies - Blue-Green



# Deployment strategies - Canary Release





Helm

Monitoring

Metrics

Utilization  
Saturation  
Errors

Application X

Application Y

Application Z

Cluster resources

Monitoring

Metrics

Rate  
Errors  
Duration

Application X

Application Y


Application Z

Cluster resources




# Patterns


# Single Concern Principle




# High Observability Principle




# Life-cycle Conformance Principle




# Process Disposability Principle




# Image Immutability Principle



# Self-containment Principle




# Sidecar Pattern







# Adapter Pattern



# Ambassador Pattern



# Initializer Pattern



# Self-awareness Pattern

